

# PARALLEL PROGRAMS, PROCESSES AND THREADS



# AGENDA

- The problem - A Case
- A Solution - Parallelism
- Processes and Threads in Linux
- Advantages & Disadvantages with multitasking

# CASE



AARHUS  
UNIVERSITY

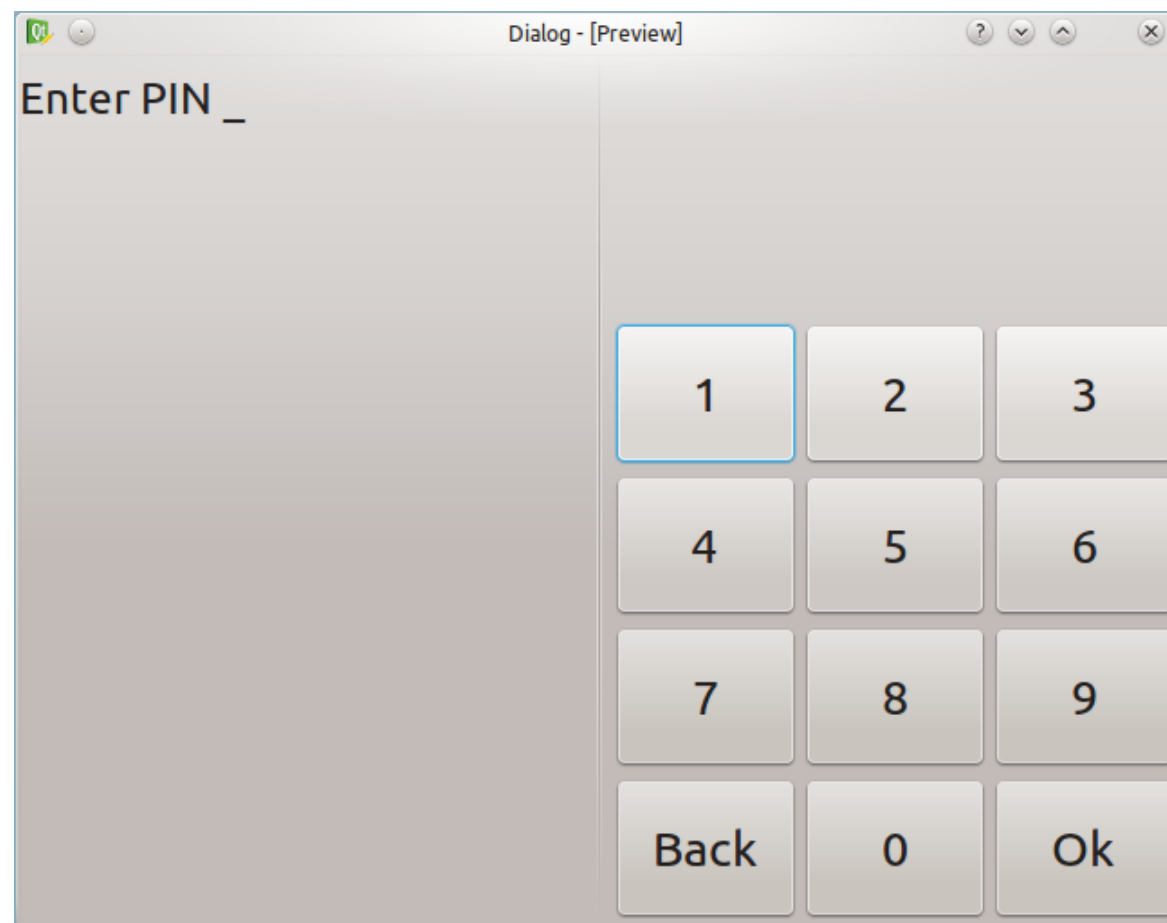
AARHUS UNIVERSITY SCHOOL OF ENGINEERING



# CASE

- Consider a system that allows a user to enter a PIN.
  - *How would you implement this?*

## BANK TERMINAL



Dialog - [Preview]

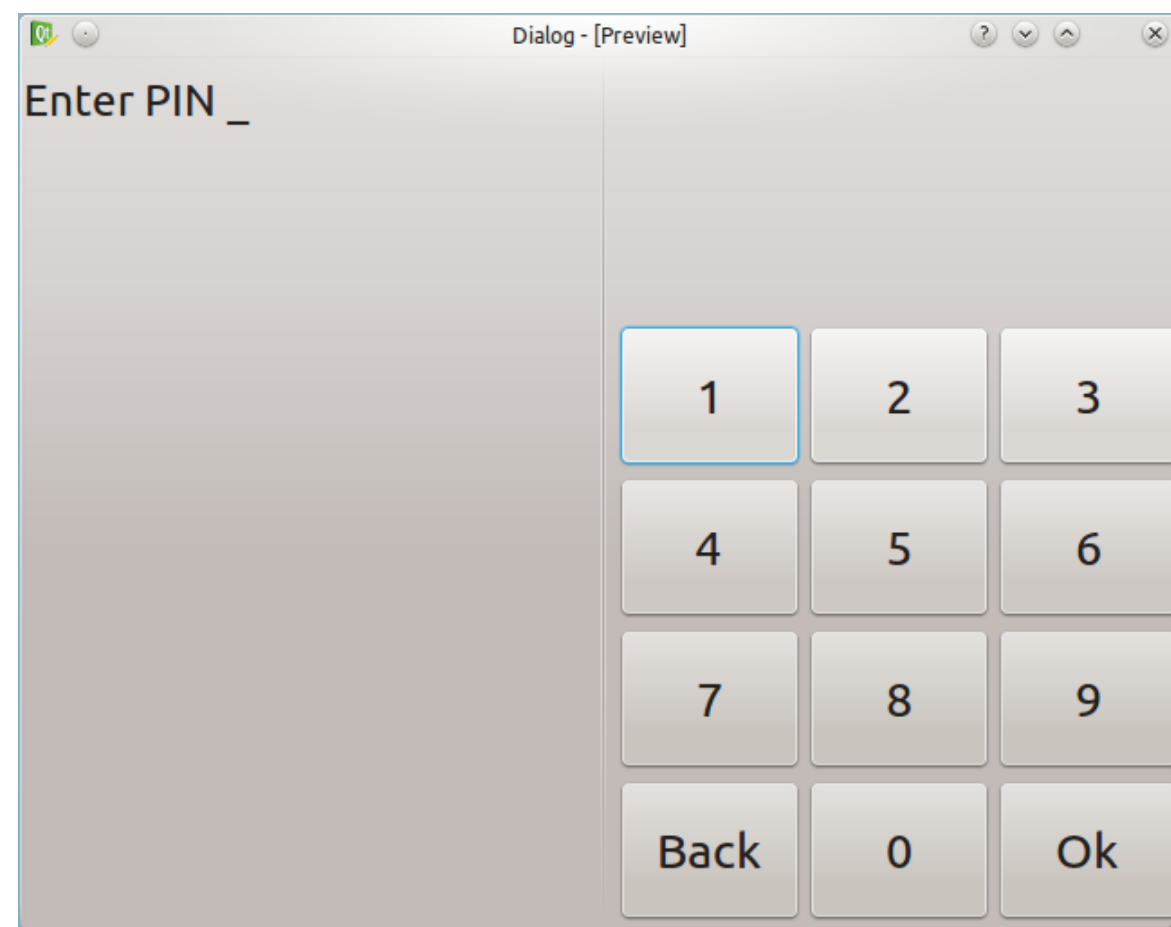
Enter PIN \_

1	2	3
4	5	6
7	8	9
Back	0	Ok

# CASE

- Consider a system that allows a user to enter a PIN.
  - *How would you implement this?*

## BANK TERMINAL




## POSSIBLE IMPLEMENTATION

```
01 void main()
02 {
03     print("ENTER PIN:");
04     ch = getKey();
05     while(ch != "OK")
06     {
07         input += ch
08         compare(input, pin)
09         ...
10 }
```

# CASE

- Now consider the same system, but now with a clock.
  - *How would you implement this?*

## BANK TERMINAL




- How would you update the clock while waiting for input?
- How would you capture key presses while updating the clock?

# CASE

- Now consider the same system, but now with a clock.
  - *How would you implement this?*

## BANK TERMINAL



## POSSIBLE IMPLEMENTATION

```
01 void main()  
02 {  
03     ???  
04 }
```

- How would you update the clock while waiting for input?
- How would you capture key presses while updating the clock?

# CASE - SOLUTION

## BANK TERMINAL

Dialog - [Preview]

Enter PIN \_


12:35:03

1	2	3
4	5	6
7	8	9
Back	0	Ok



# CASE - SOLUTION

## BANK TERMINAL




## USER INPUT THREAD

- Waiting on blocking call

```
01 void userInput()  
02 {  
03     print("ENTER PIN:");  
04     ch = getKey();  
05     while(ch != "OK")  
06     {  
07         input += ch  
08         compare(input, pin)  
09         ...  
10     }  
11 }
```

# CASE - SOLUTION

## BANK TERMINAL



## USER INPUT THREAD

- Waiting on blocking call

```
01 void userInput()  
02 {  
03     print("ENTER PIN:");  
04     ch = getKey();  
05     while(ch != "OK")  
06     {  
07         input += ch  
08         compare(input, pin)  
09         ...  
10     }  
11 }
```

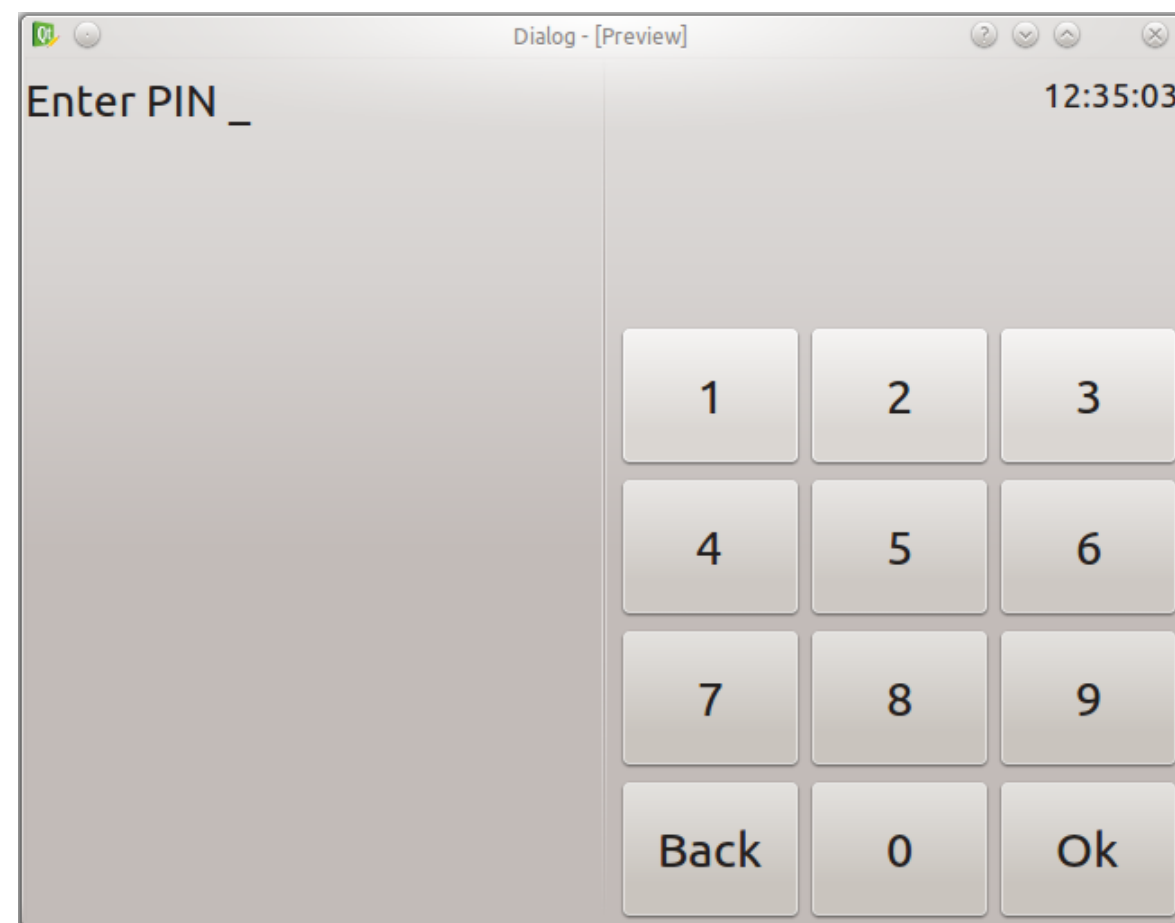
## CLOCK/TIME THREAD

- Delay is blocking

```
01 void updClock()  
02 {  
03     while(true)  
04     {  
05         display current time  
06         wait 1s  
07     }  
08 }
```

# CASE - SOLUTION

## BANK TERMINAL



## USER INPUT THREAD

- Waiting on blocking call

```
01 void userInput()  
02 {  
03     print("ENTER PIN:");  
04     ch = getKey();  
05     while(ch != "OK")  
06     {  
07         input += ch  
08         compare(input, pin)  
09         ...  
10     }  
11 }
```

## CLOCK/TIME THREAD

- Delay is blocking

```
01 void updClock()  
02 {  
03     while(true)  
04     {  
05         display current time  
06         wait 1s  
07     }  
08 }
```

## STARTS THREADS

```
01 int main()  
02 {  
03     startThread(userInput)  
04     startThread(updClock)  
05 }
```

# PARALLEL PROGRAMS - A SOLUTION

- *Parallel programs* are programs
  - where the work done is parallelised
- Parallelisation can take the form of
  - Multiple processes working in conjunction
  - Multiple threads, jobs or tasks in process working in conjunction



# PROCESS & THREADS IN LINUX



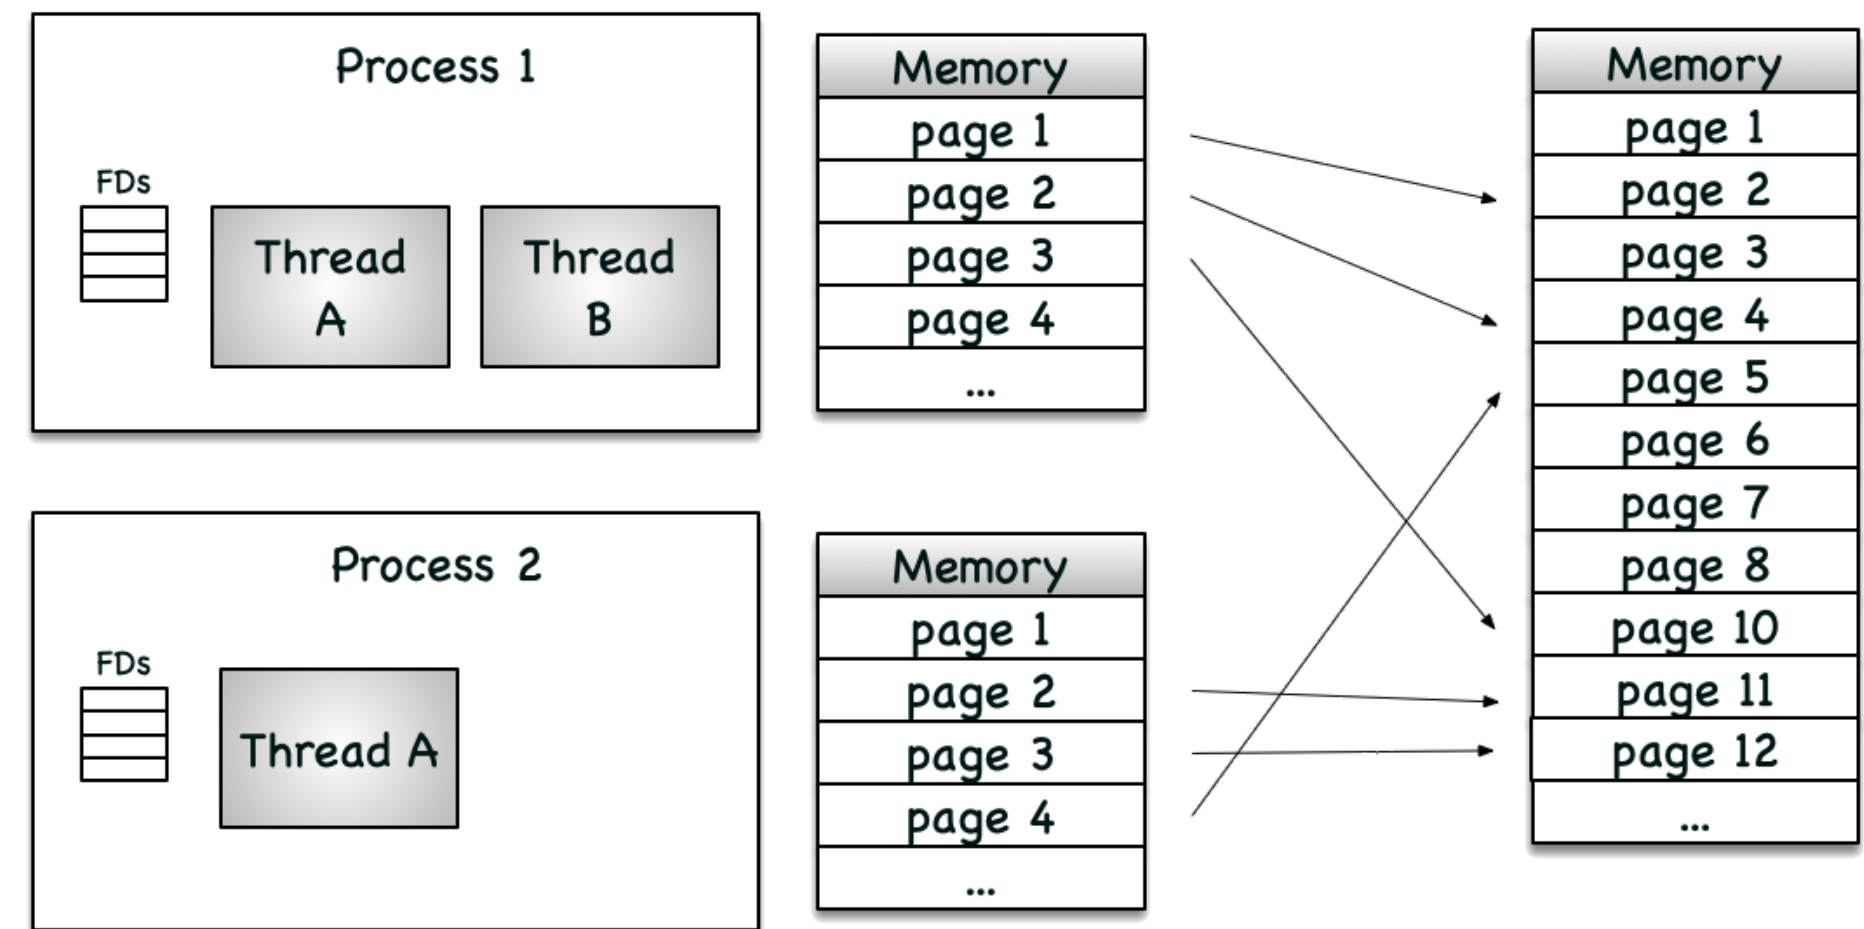
# PROCESS & THREADS IN LINUX

- Processes and the OS (Kernel)
- Process anatomy
- Shared memory



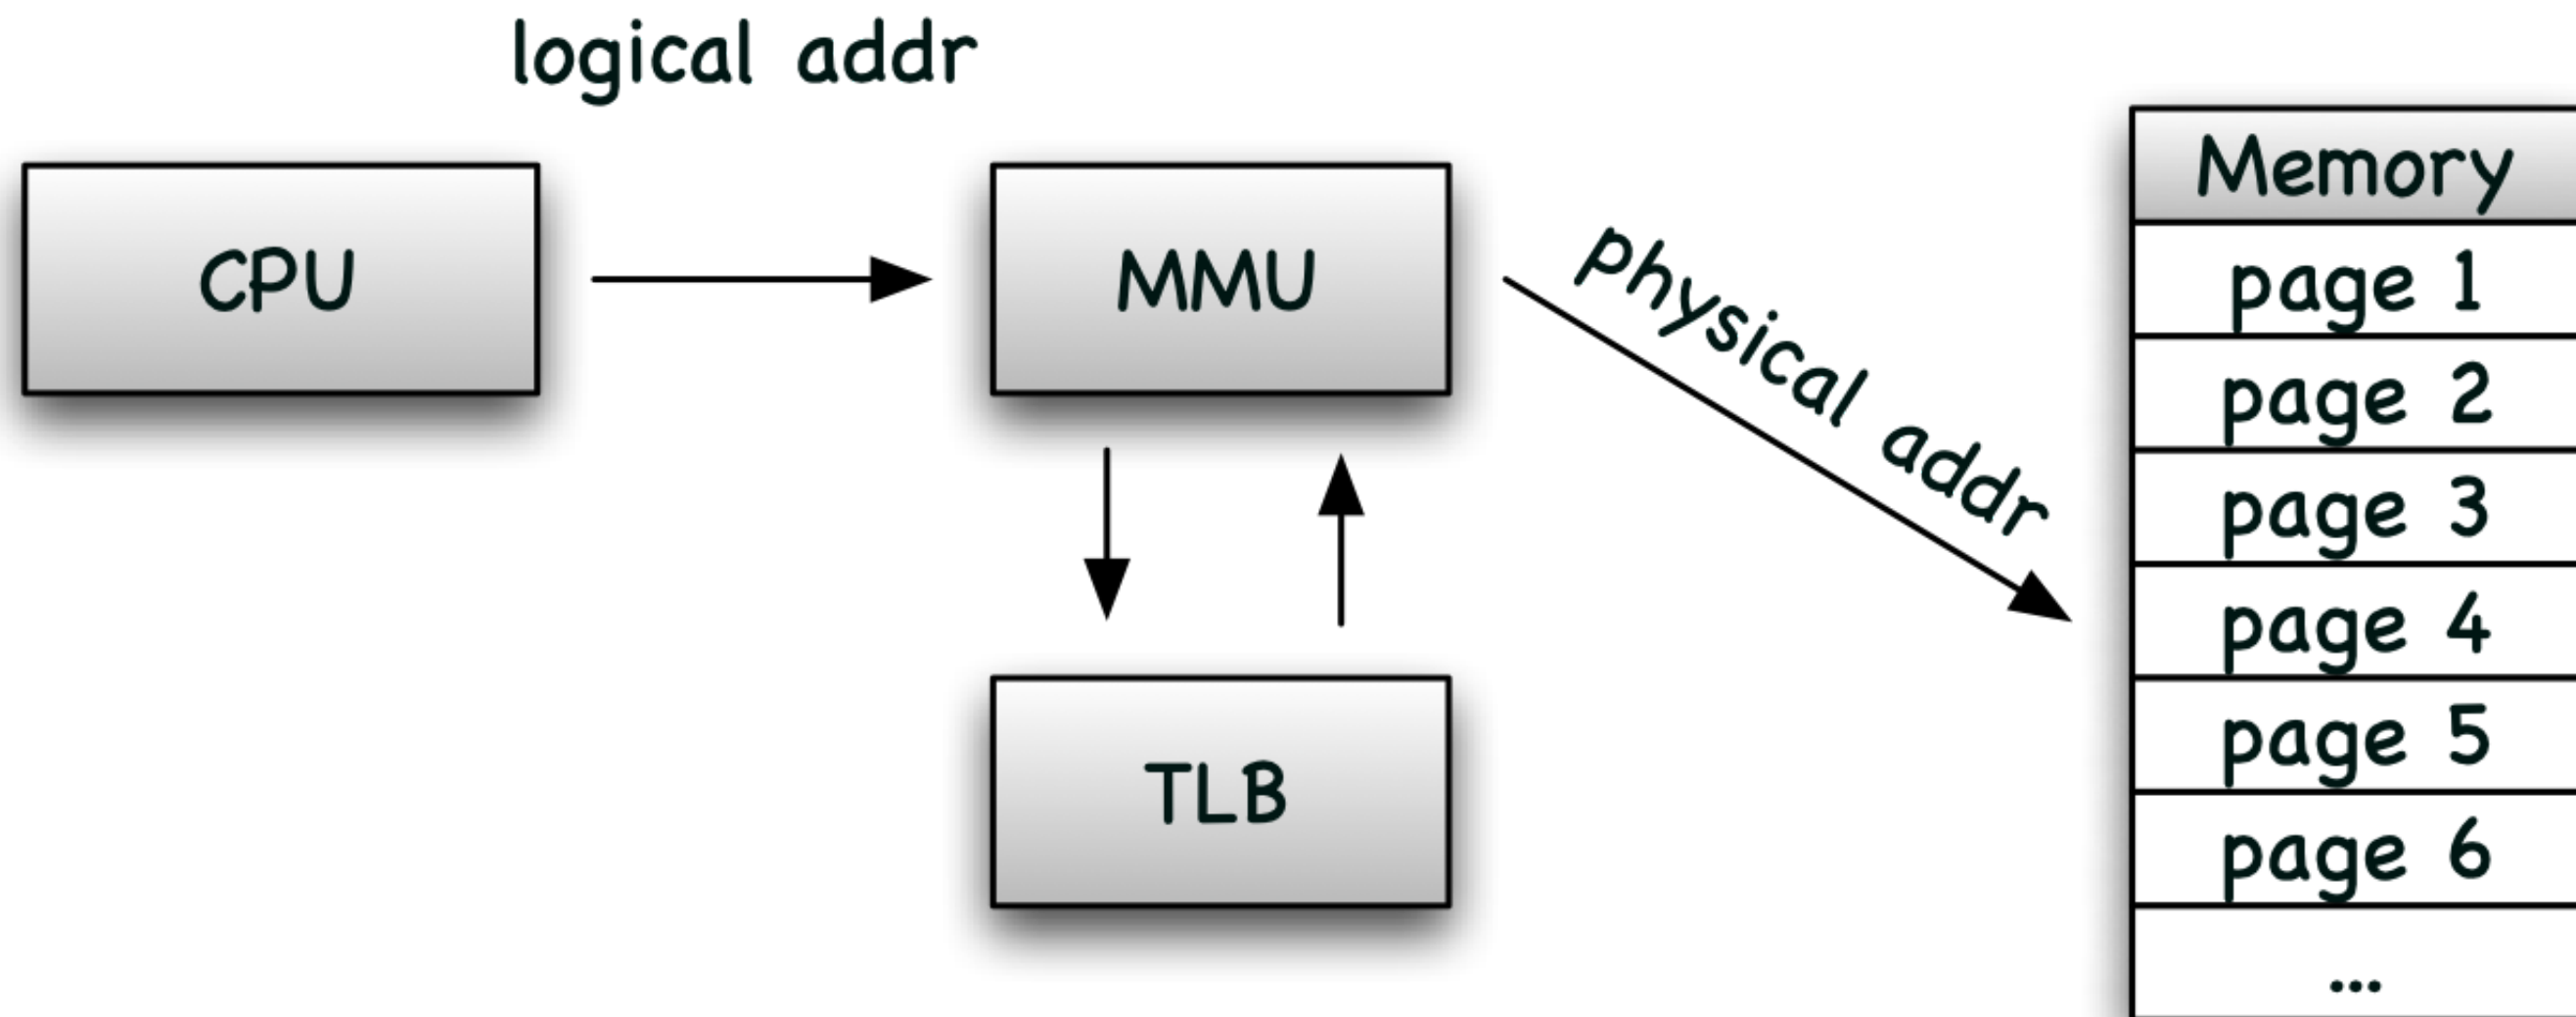
# PROCESSES AND THE OS (KERNEL)

- Threads, Processes and Memory mapping
  - Each process has its own memory space
    - A mapping exists between virtual and physical memory
    - Not possible for one process to write in another address space
- Threads share data space
  - Care must be taken NOT to destroy other threads data



# PROCESSES AND THE OS (KERNEL)

## MMU - MEMORY MANAGEMENT UNIT



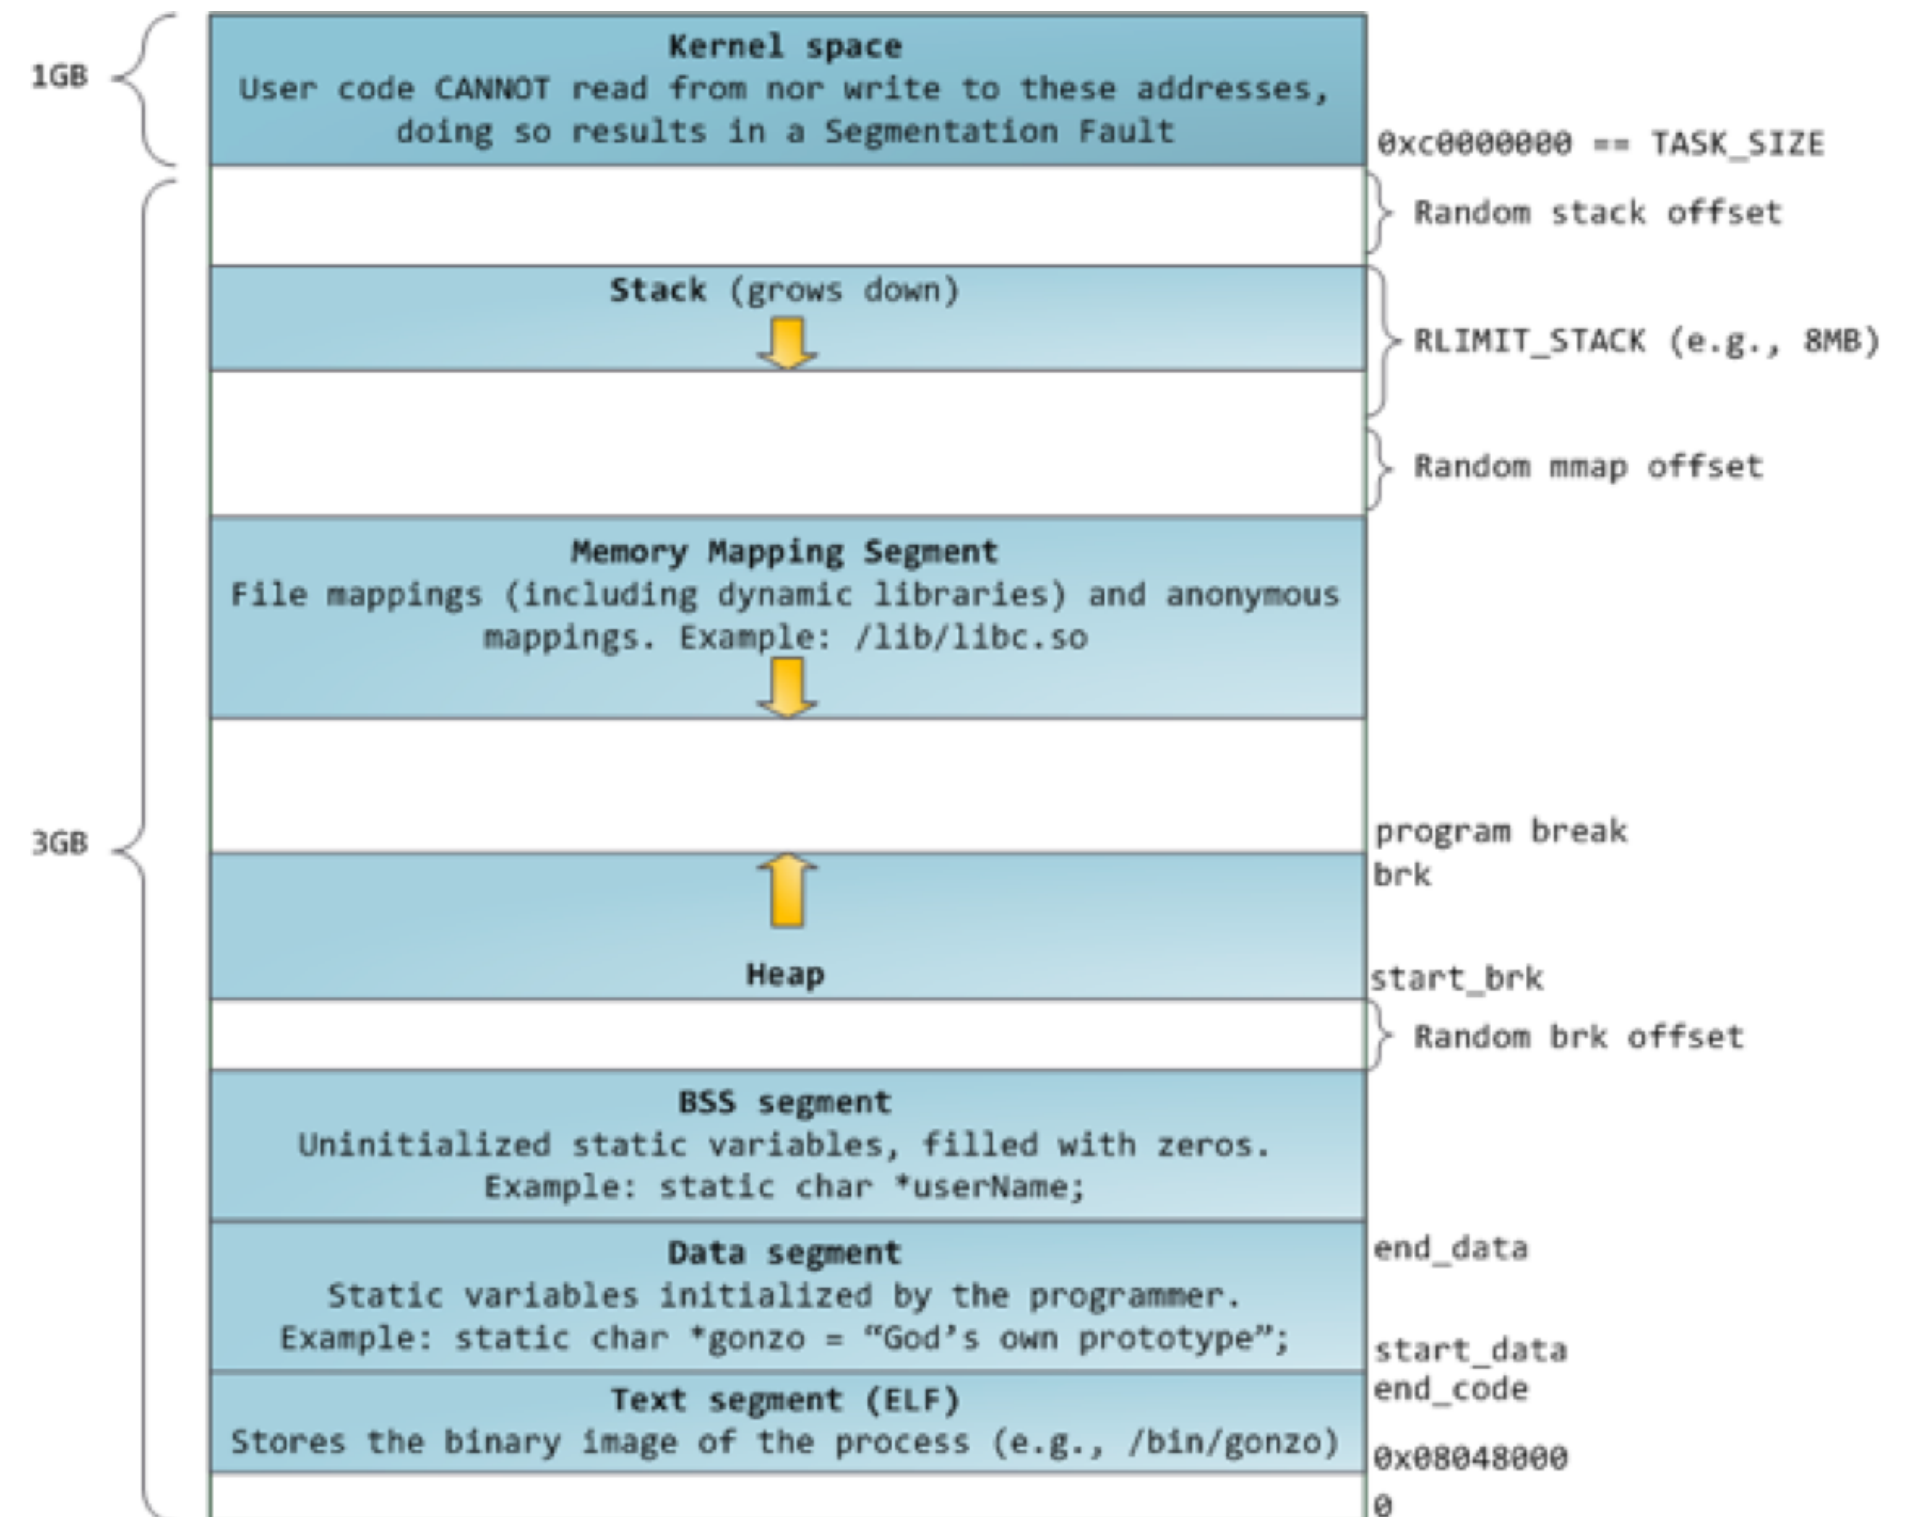
MMU: Memory Management Unit  
TLB: Translation Lookaside Buffer



# PROCESS ANATOMY

## PROCESS - A PROGRAM BEING EXECUTED IN LINUX

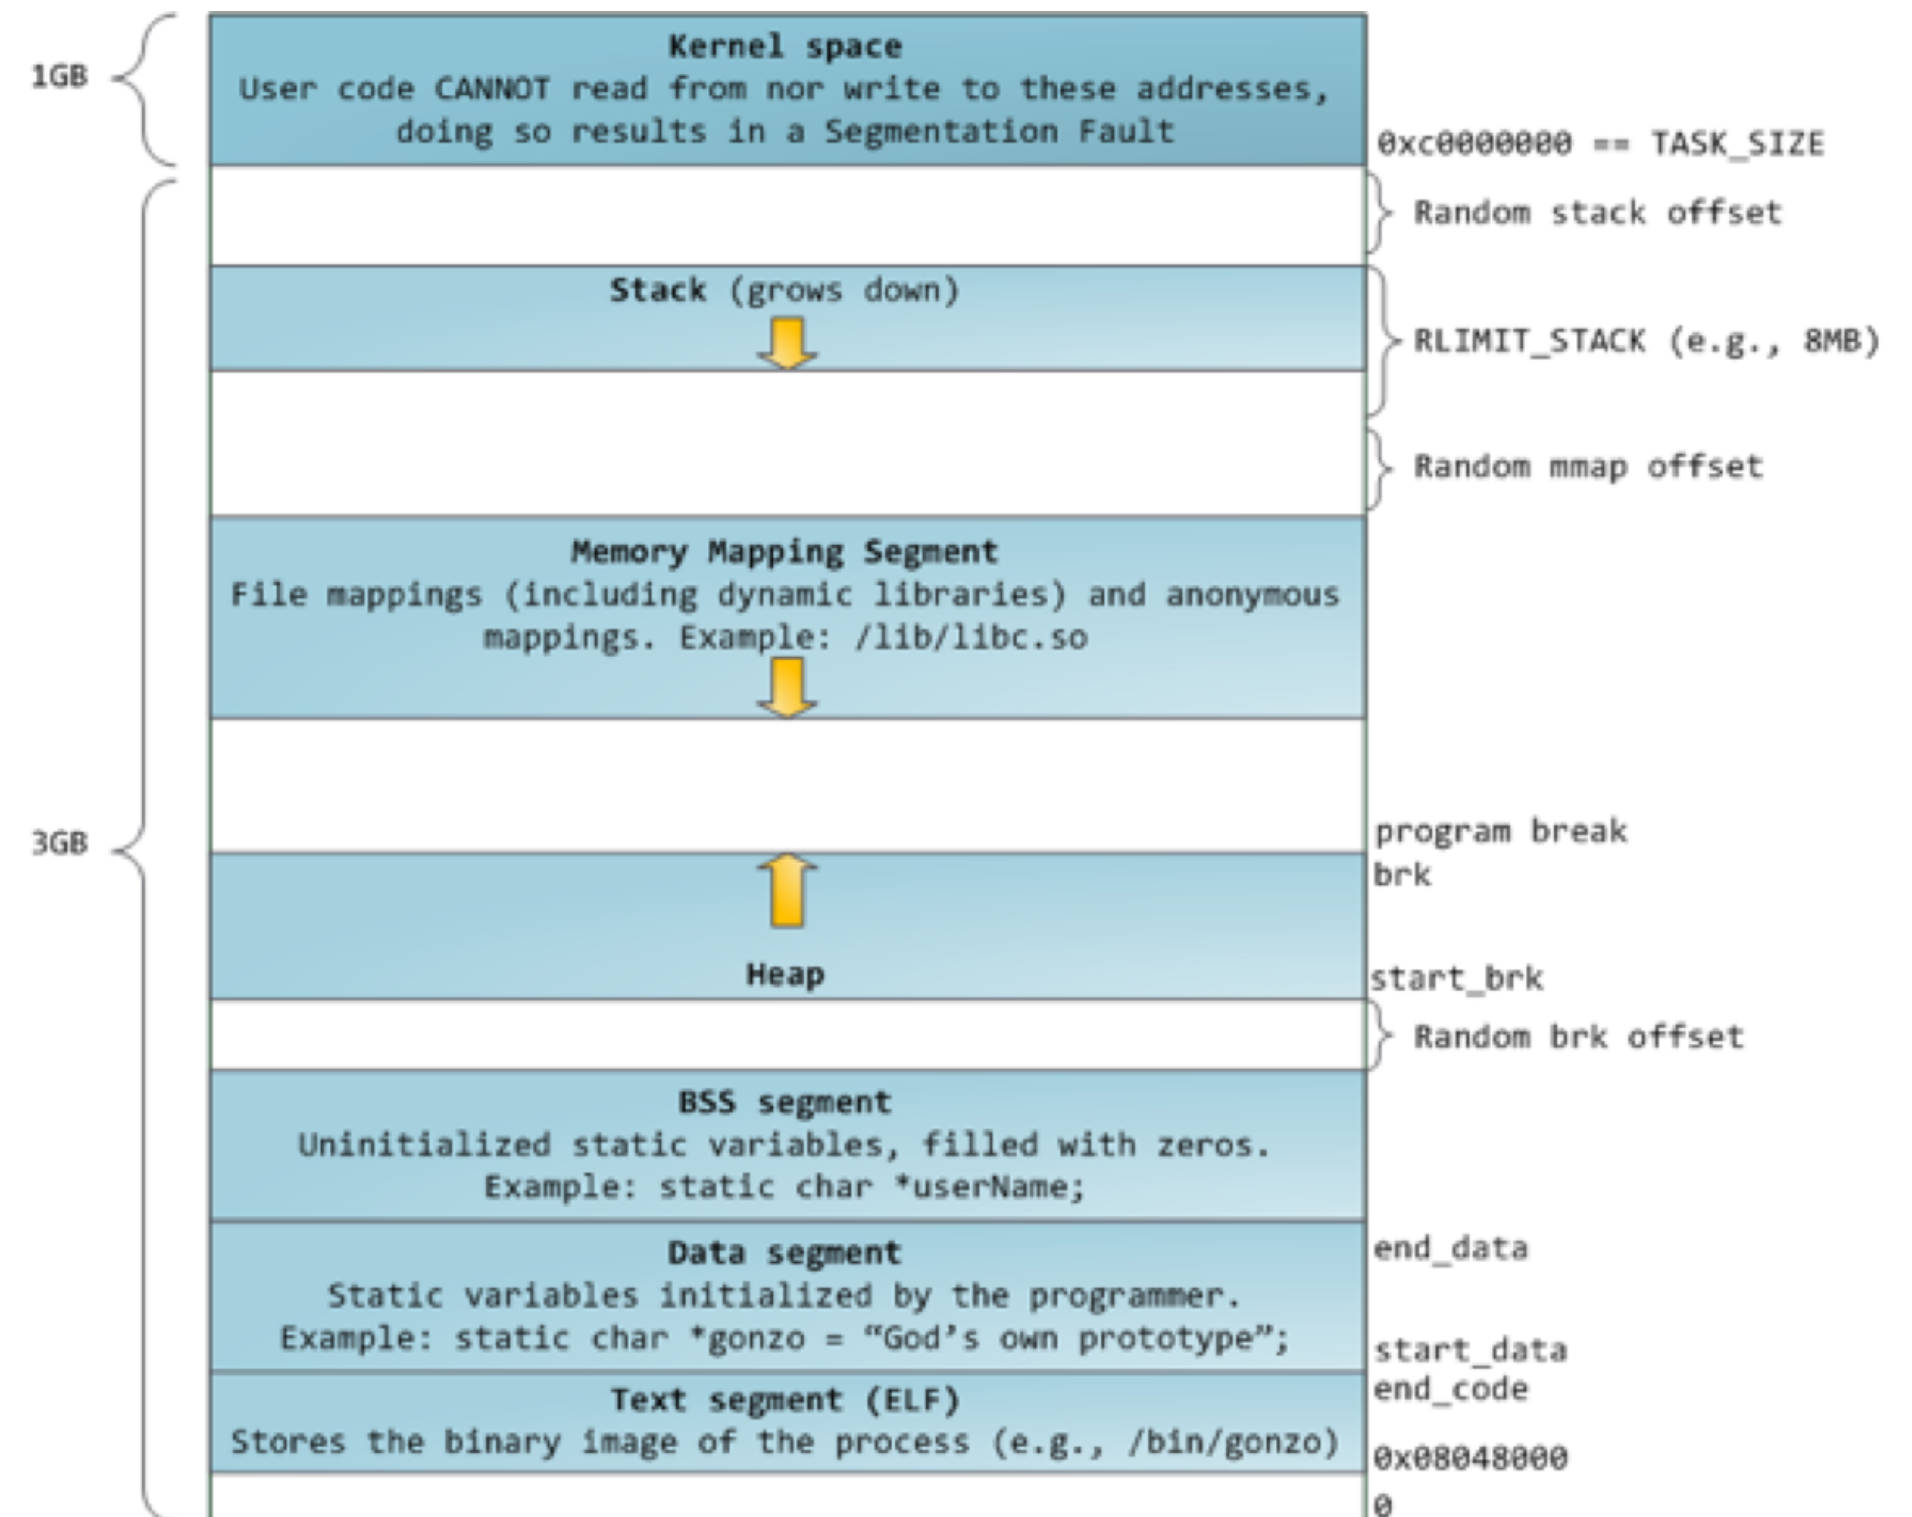
- Stack
  - Local variables
  - Function return values
  - LIFO



# PROCESS ANATOMY

## PROCESS - A PROGRAM BEING EXECUTED IN LINUX

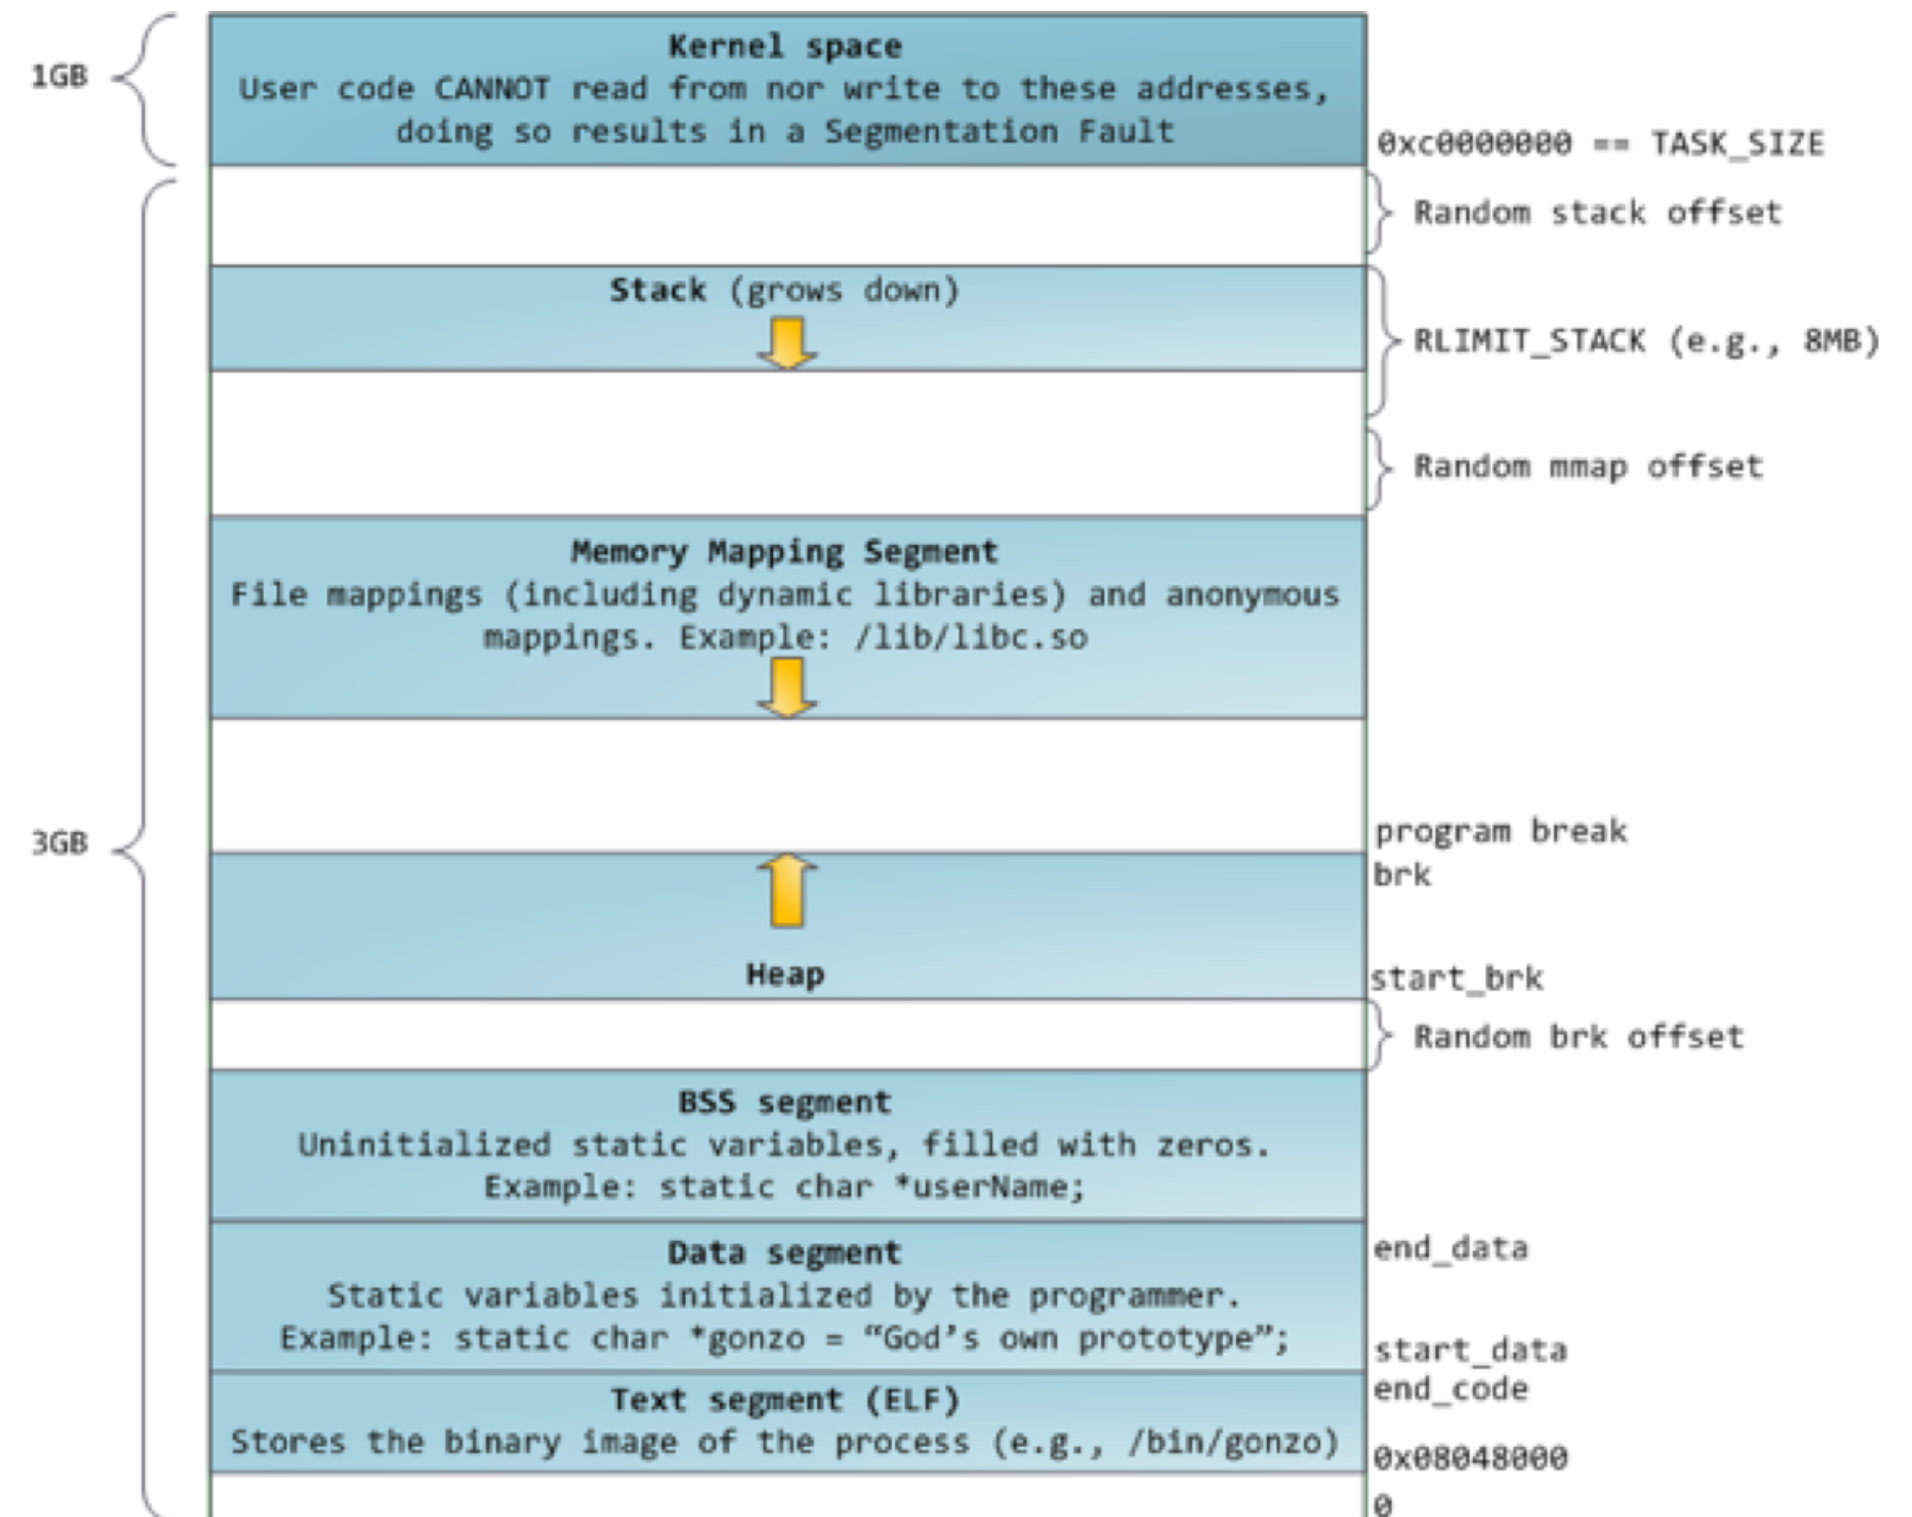
- Stack
  - Local variables
  - Function return values
  - LIFO
- Heap
  - "Free-store"
  - Dynamically allocated memory



# PROCESS ANATOMY

## PROCESS - A PROGRAM BEING EXECUTED IN LINUX

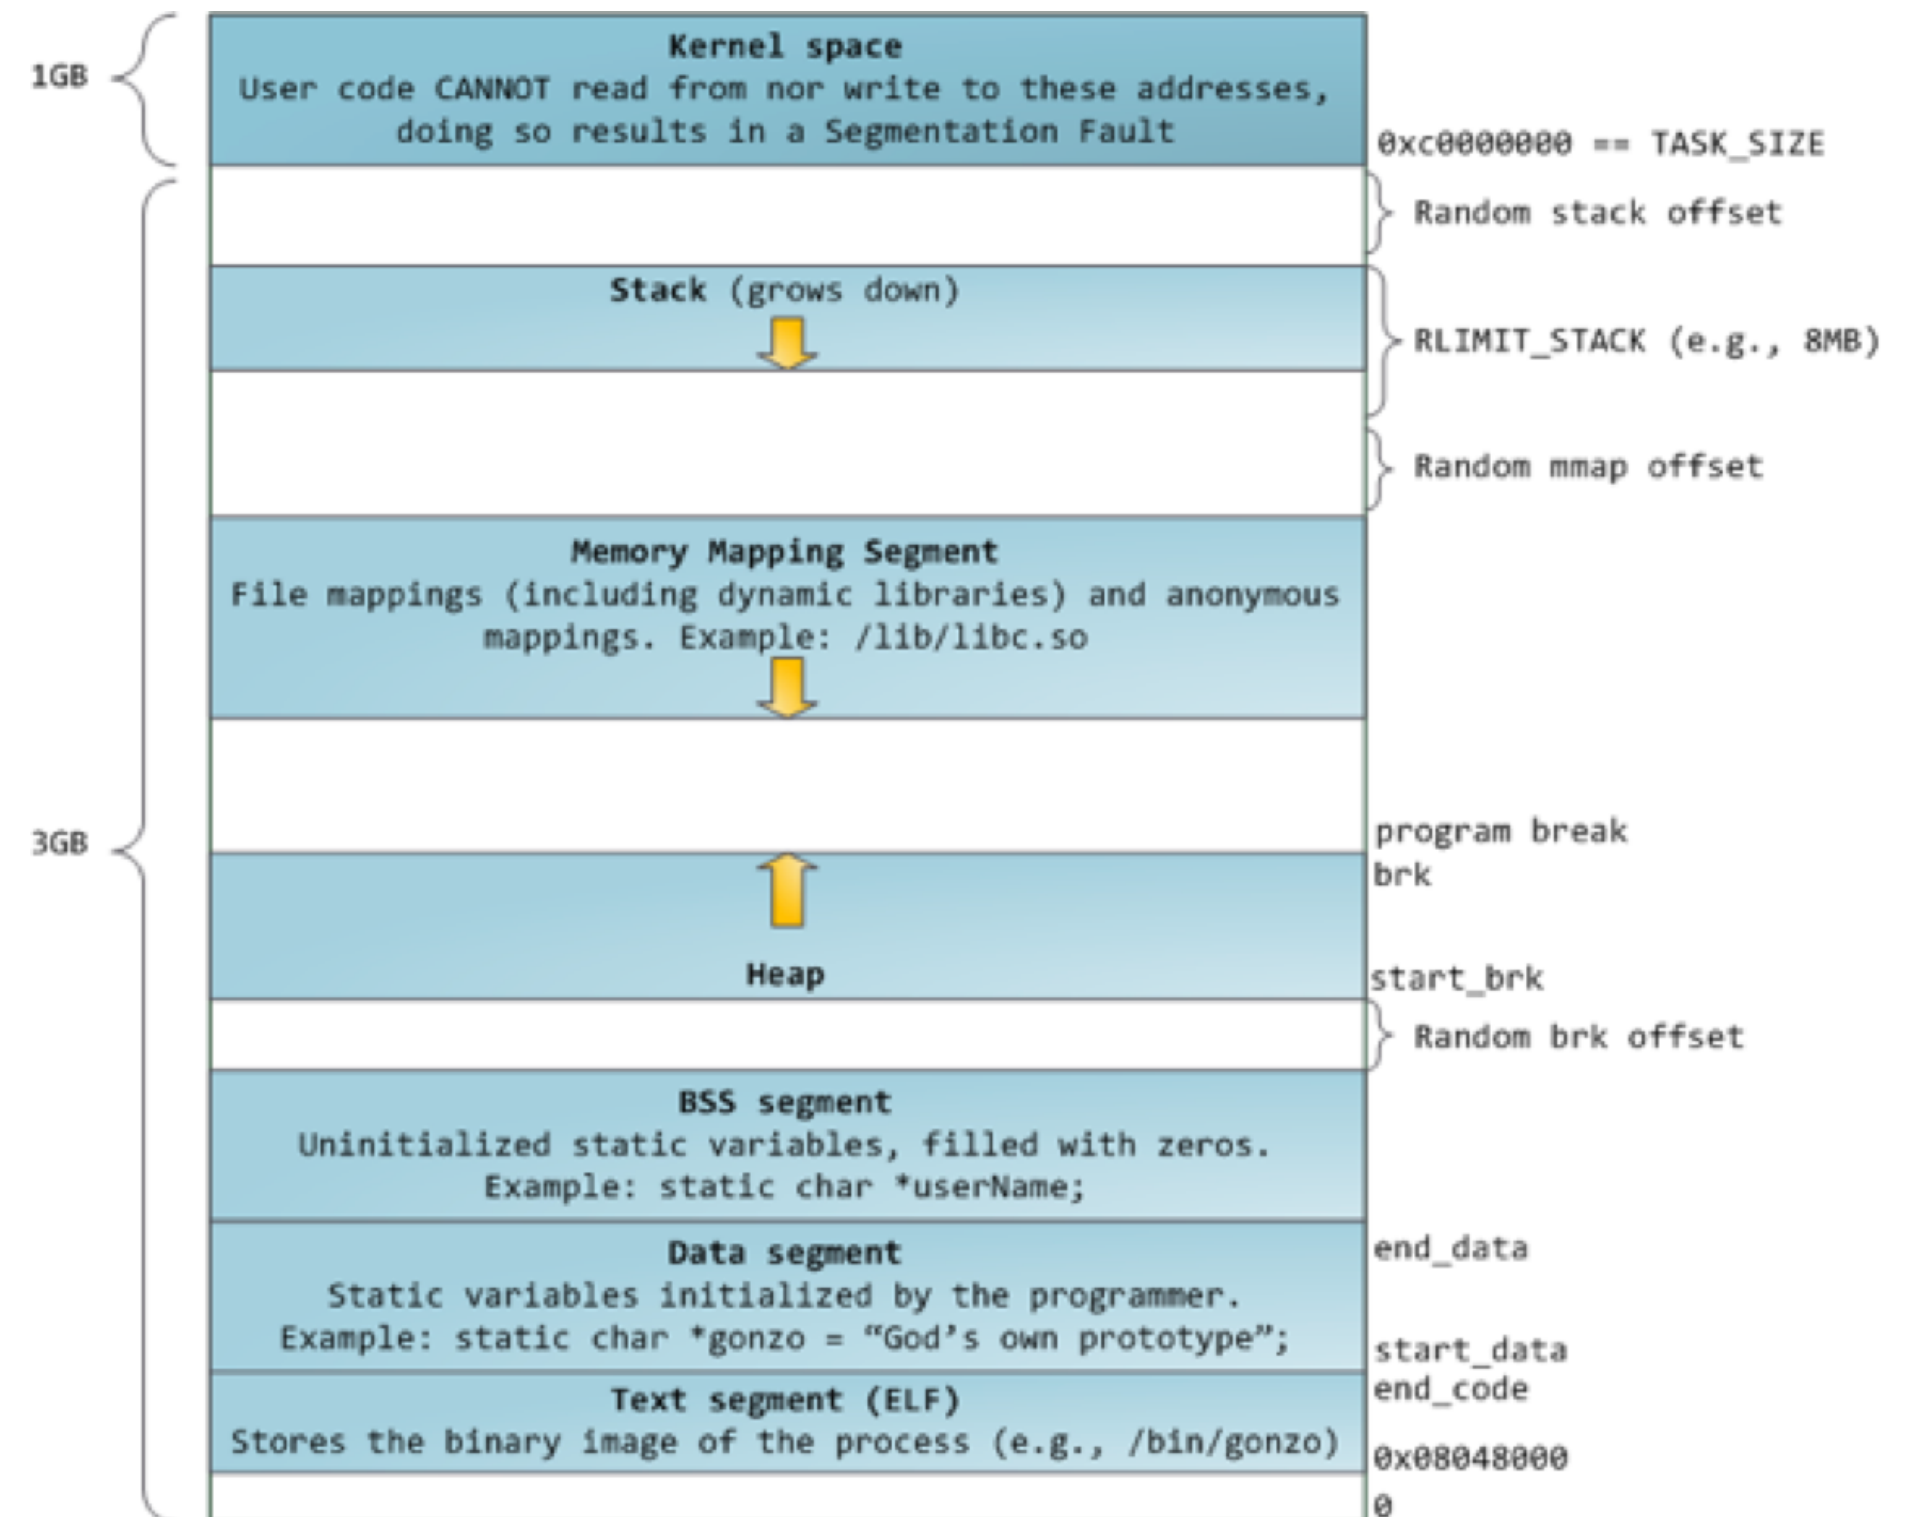
- Stack
  - Local variables
  - Function return values
  - LIFO
- Heap
  - "Free-store"
  - Dynamically allocated memory
- Memory Mapping
  - File mapped in memory
  - Includes dyn libs
  - Sharing memory between processes



# PROCESS ANATOMY

## PROCESS - A PROGRAM BEING EXECUTED IN LINUX

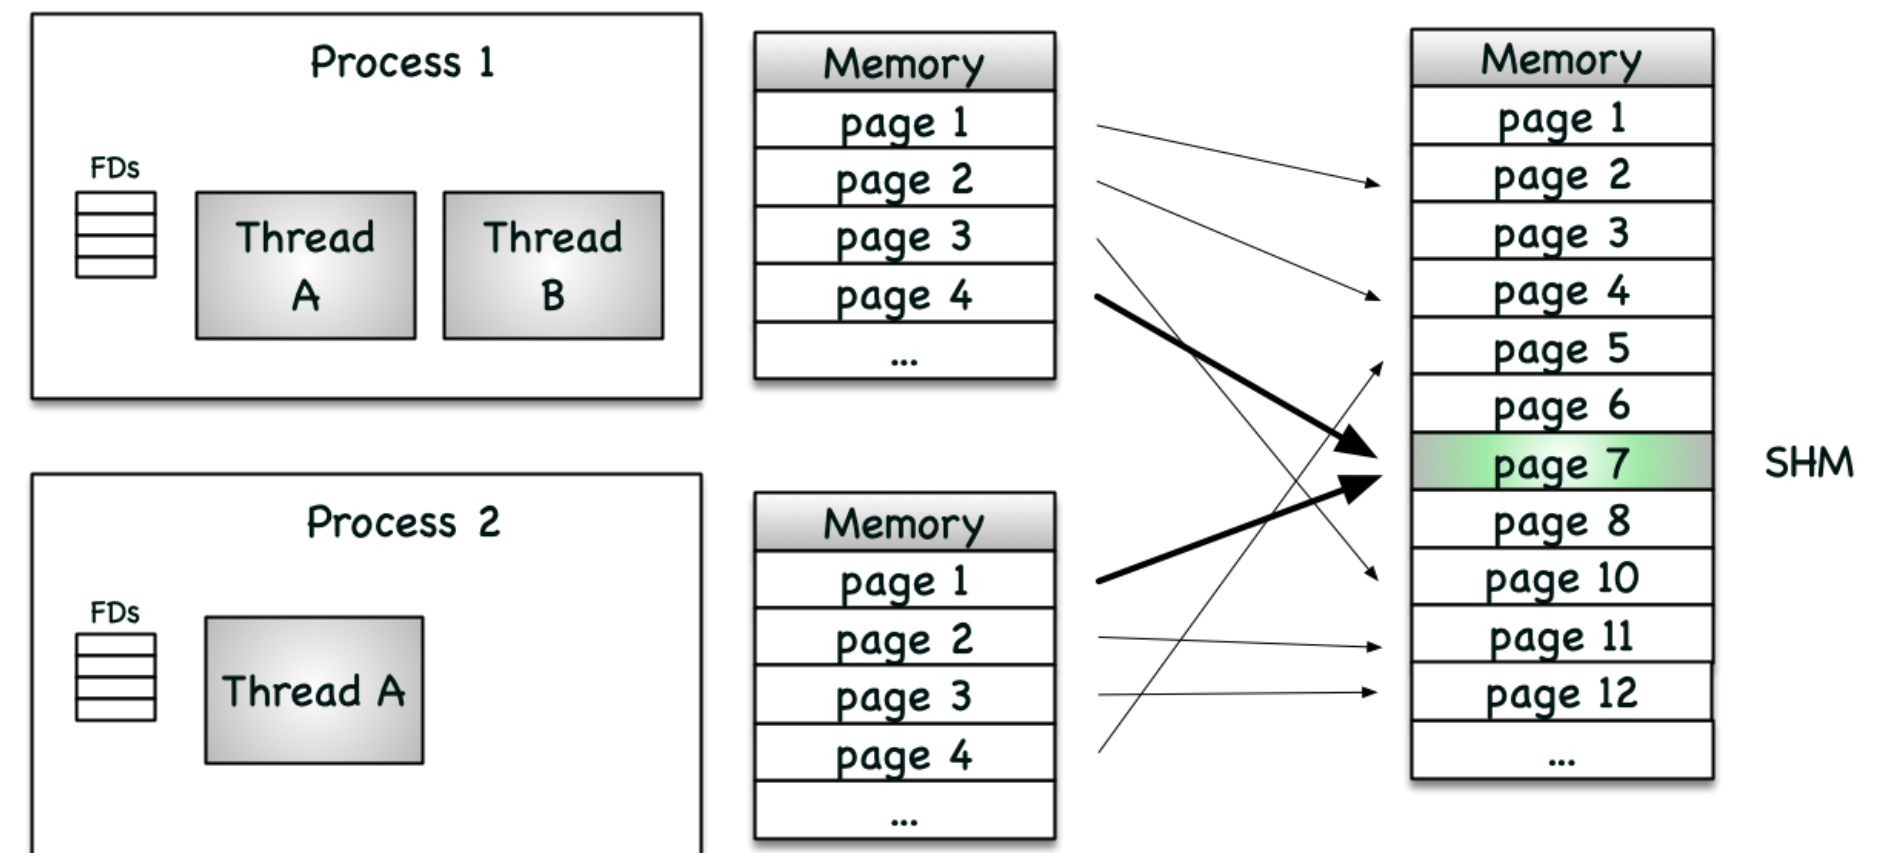
- Stack
  - Local variables
  - Function return values
  - LIFO
- Heap
  - "Free-store"
  - Dynamically allocated memory
- Memory Mapping
  - File mapped in memory
  - Includes dyn libs
  - Sharing memory between processes
- Variables and ELF





# SHARED MEMORY

- POSIX Shared Memory
  - Accessing memory affects other process
  - Pro
    - Speed/Performance
  - Cons
    - Fragile
      - Death of process
      - Data must abide certain principles
    - Challenge ensuring synchronicity



# THREADING MODELS



AARHUS  
UNIVERSITY

AARHUS UNIVERSITY SCHOOL OF ENGINEERING

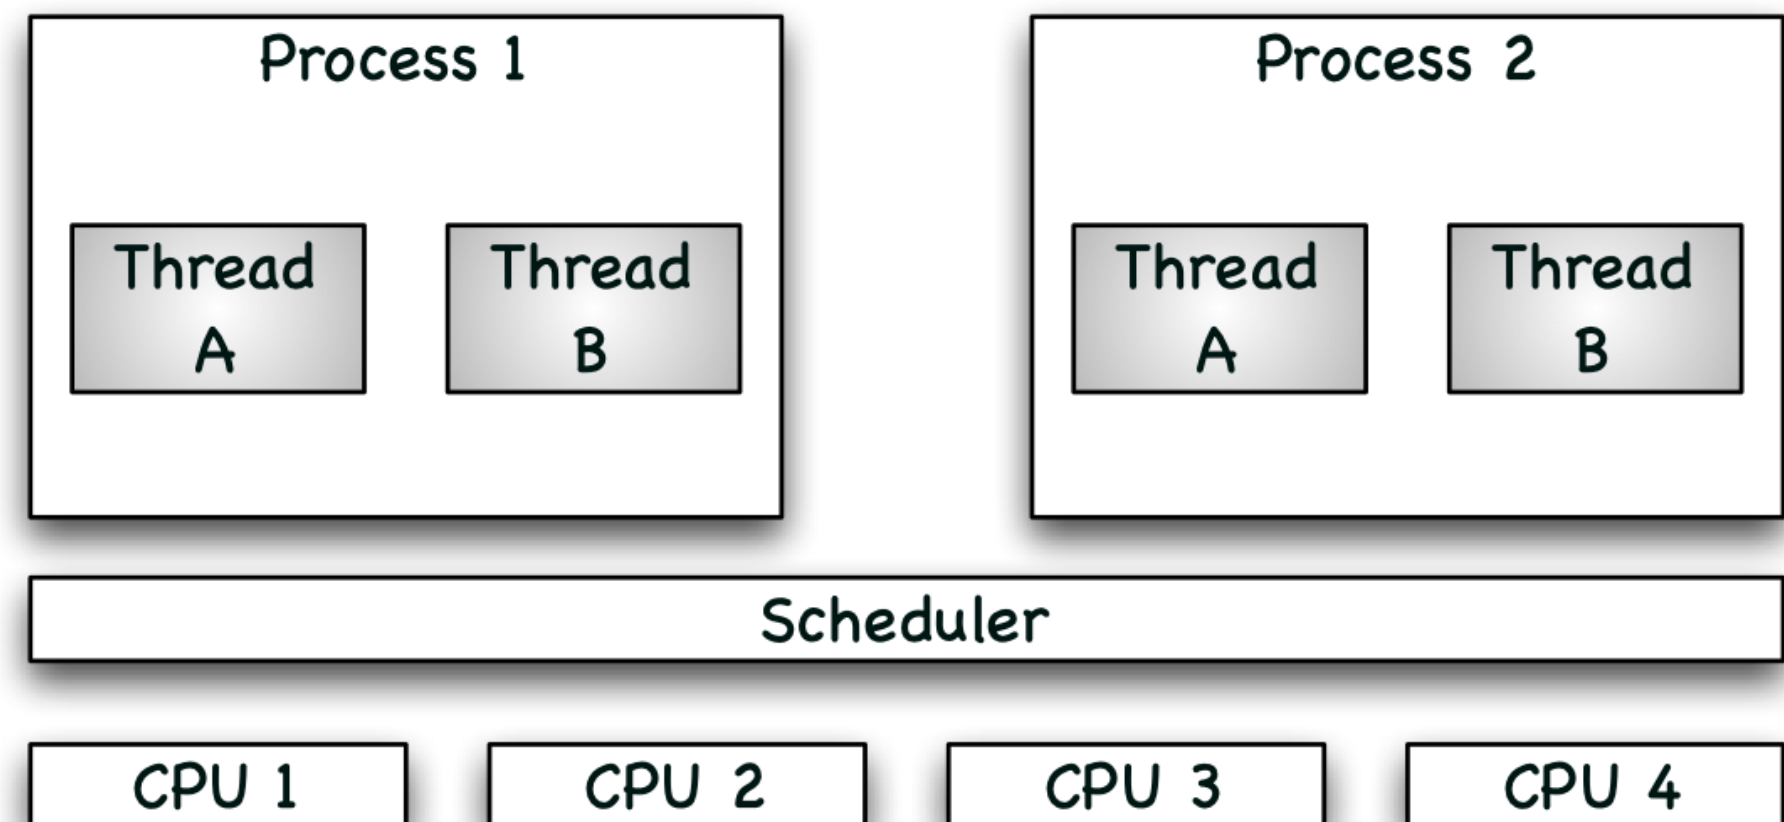


# THREADING MODELS

- Threads of execution and the OS (Kernel)
- User level threading
- Kernel level threading
- Hybrid level threading

# THREADS OF EXECUTION AND THE OS (KERNEL)


- How are threads mapped for execution?
  - By the scheduler
  - Using one of three different models
    - User level threading
    - Kernel level threading
    - Hybrid level threading






# USER-LEVEL THREADING

- Simple implementation no kernel support for threads
- Very quick thread context switch (no kernel handling needed)
- Not possible to handle multicores




# KERNEL LEVEL THREADING

- Kernel Level Threading
  - Need thread awareness in kernel
  - Maps directly to threads which the scheduler can control
  - Efficient multicore usage
- OS
  - Linux (NPLT), Win32 etc.



# HYBRID LEVEL THREADING

- Hybrid Level Threading
  - Complex implementation
  - Requires good coordination between user land and kernel land scheduler
    - Otherwise suboptimal resource usage
- OS
  - Windows 7



# CONTEXT SWITCHING



AARHUS  
UNIVERSITY

AARHUS UNIVERSITY SCHOOL OF ENGINEERING

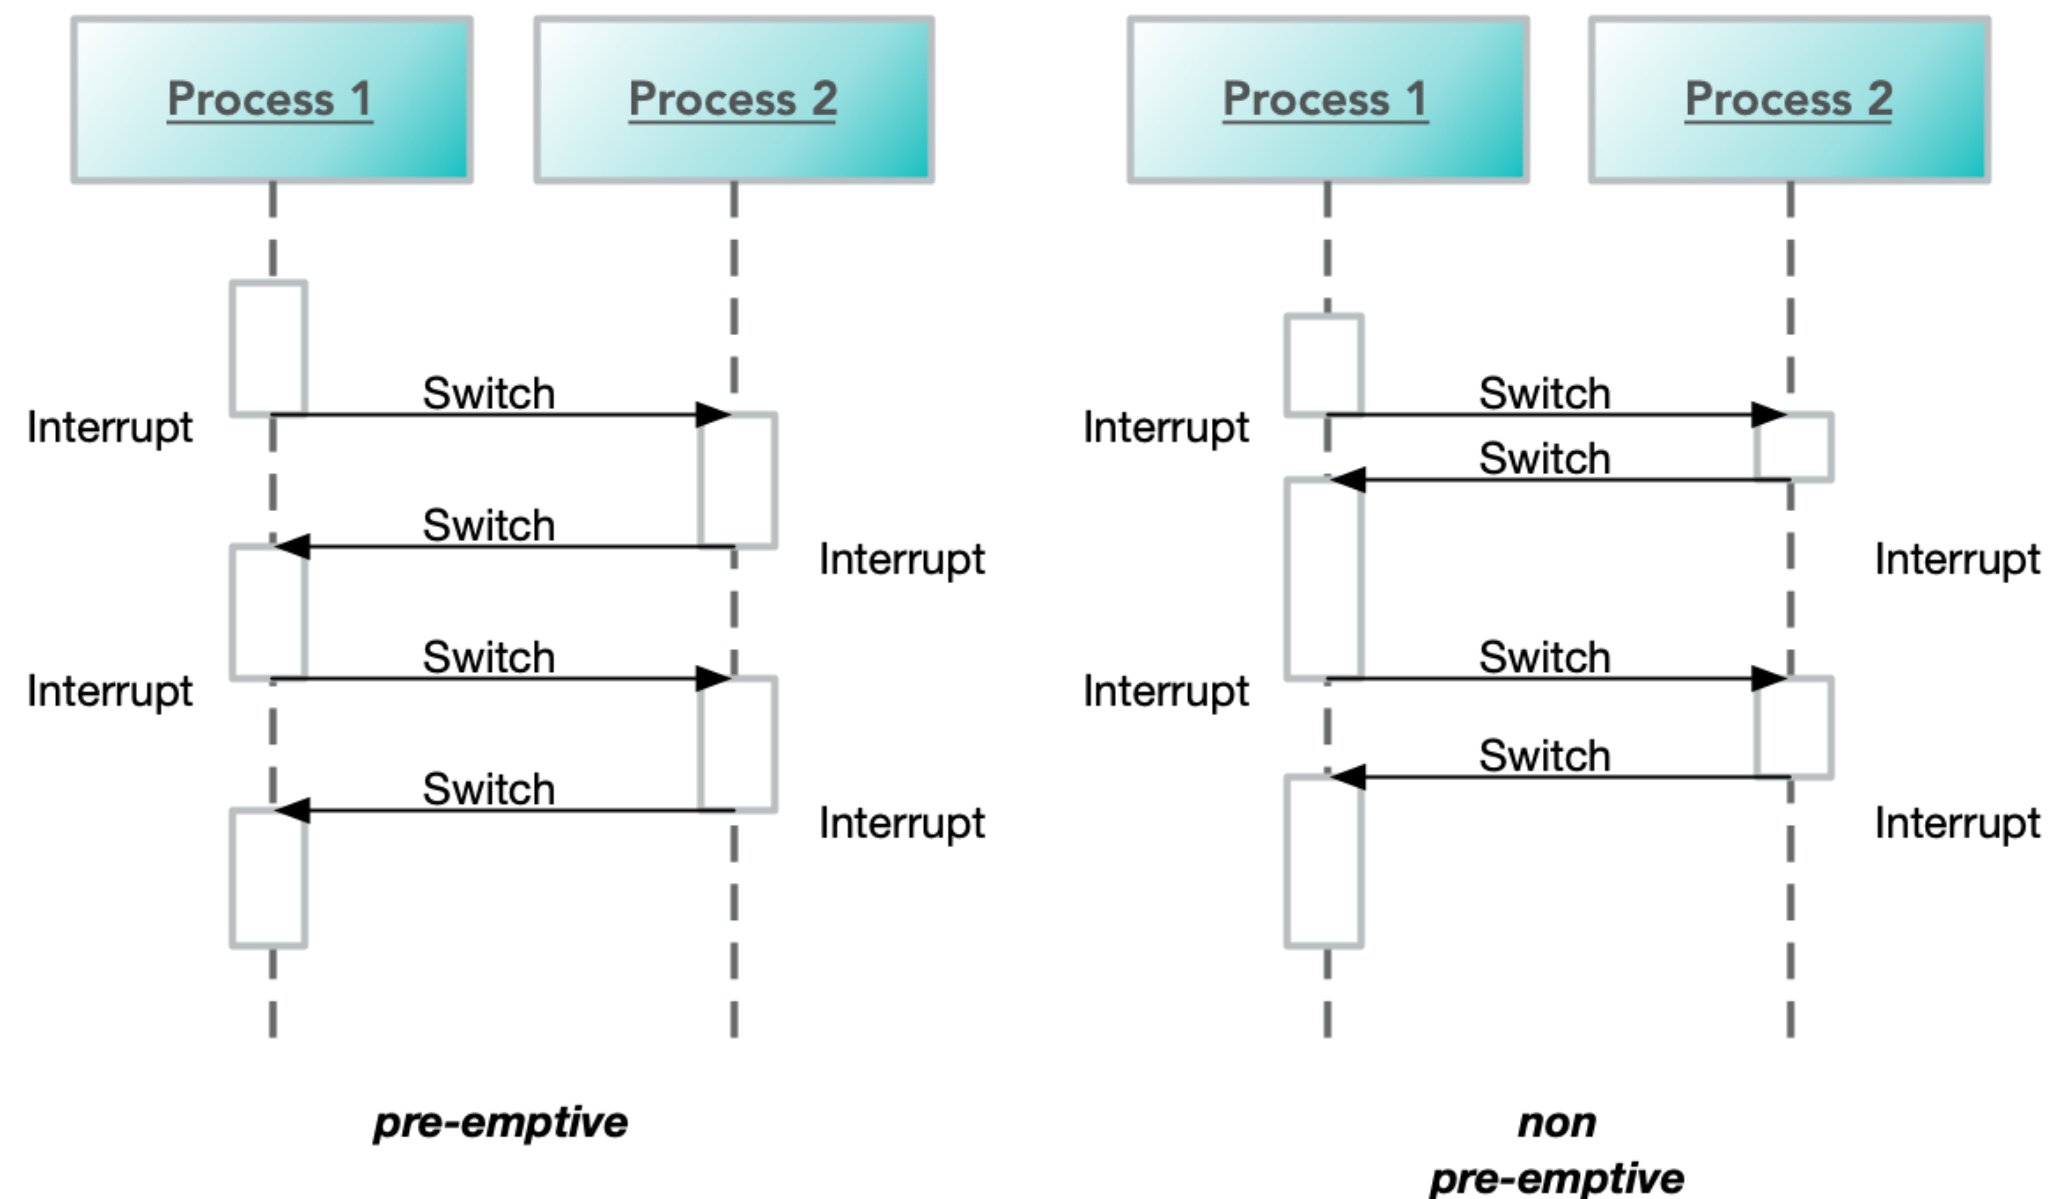


# CONTEXT SWITCHING

- A context is the environment of the currently running process
- A context switch is performed by the OS to suspend the currently running process and resume another process
- General steps:
  - Interrupt current process
  - Save context of current process (SP, PC, registers, ...)
  - Restore context of next process
  - Resume execution of next process

# CONTEXT SWITCHING

- The operating system schedules tasks for execution
  - Preemptive scheduling: Tasks can be interrupted at any time
  - Non-preemptive scheduling: Tasks voluntarily yield the CPU
- Linux supports both - Kernel configuration options

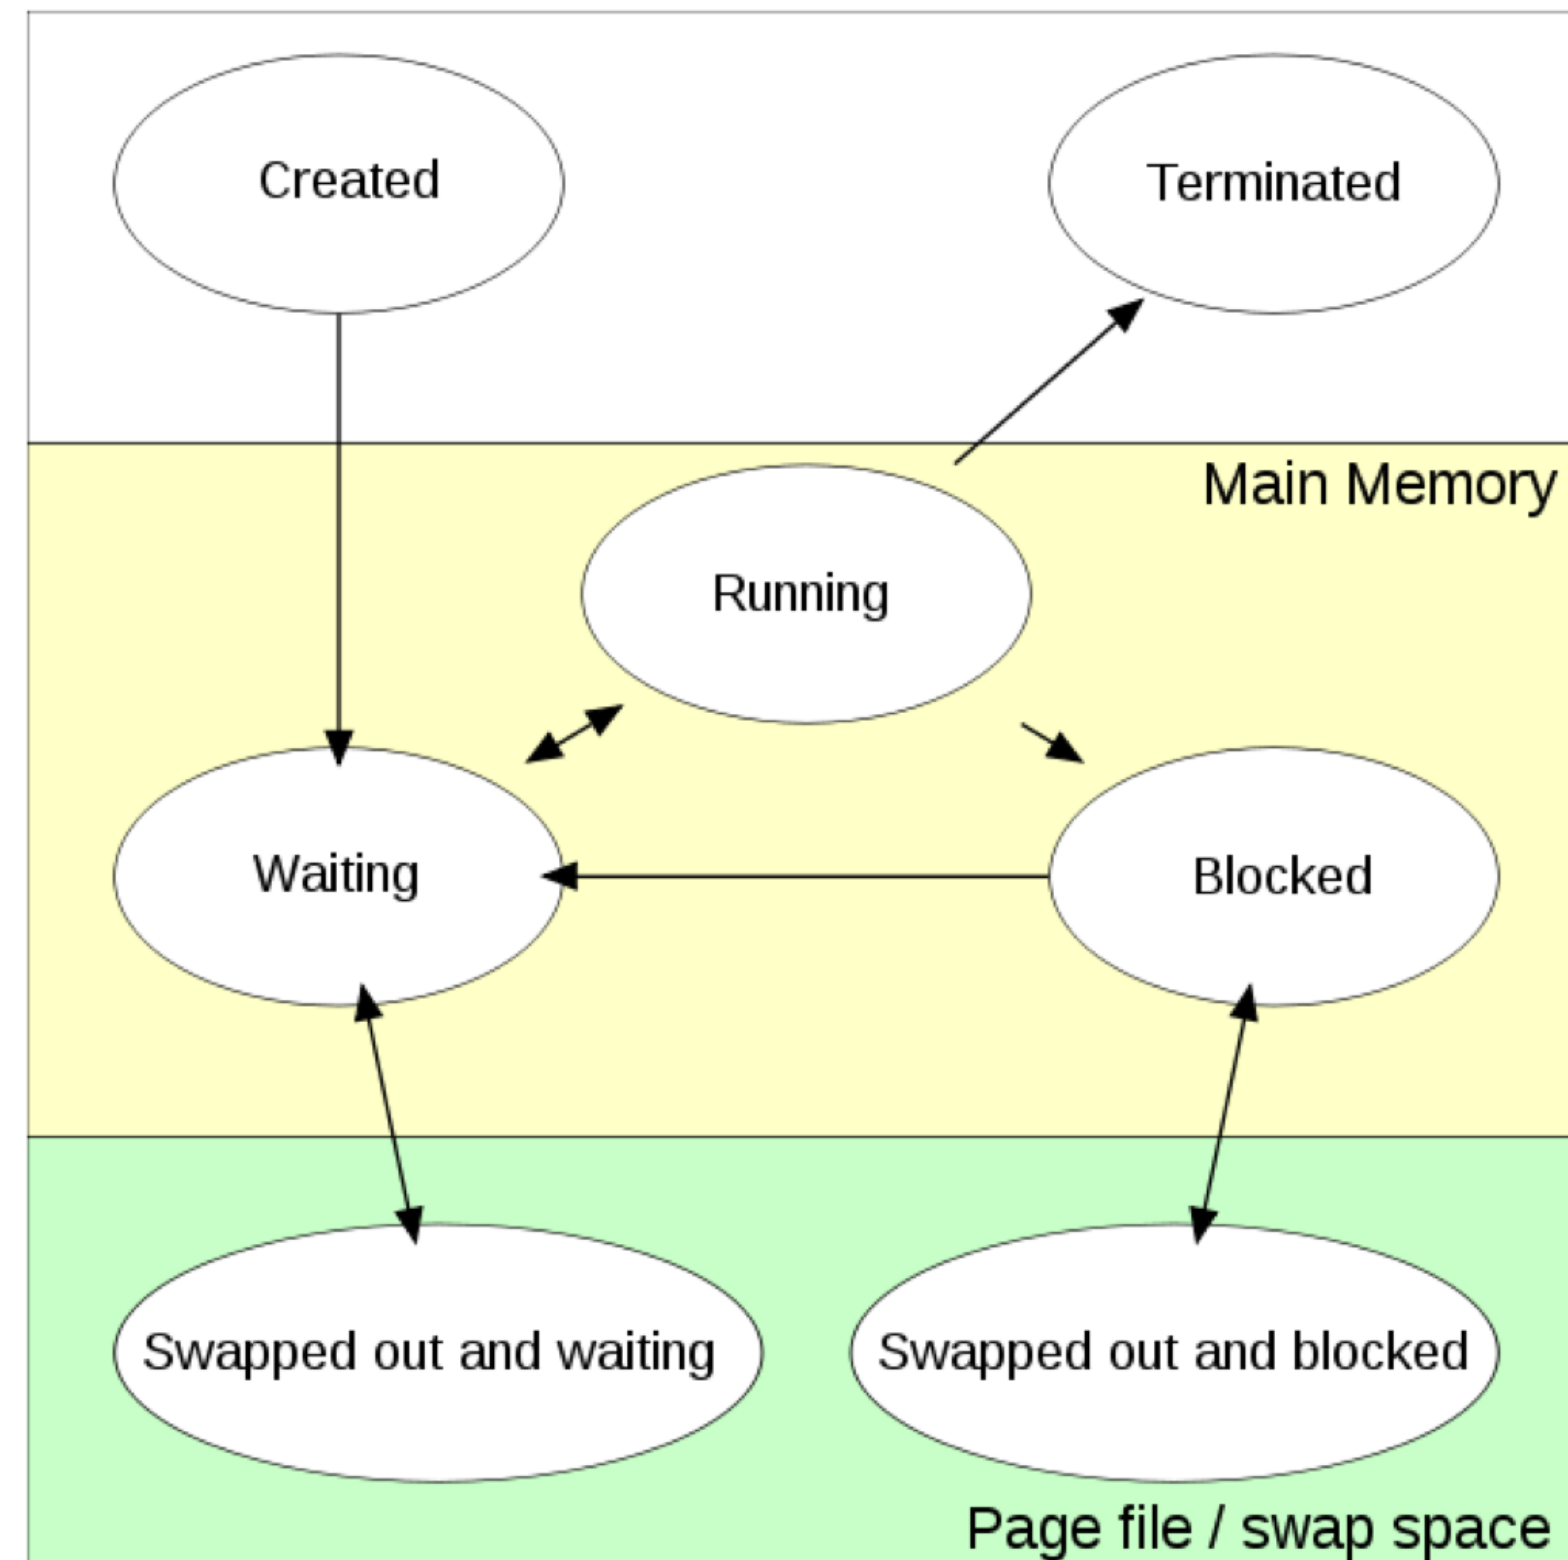


# THE LIFE AND DEATH OF A PROCESS



# THE LIFE AND DEATH OF A PROCESS

- Waiting
  - In queue to run on processor
- Running
  - Running on the processor
- Blocked
  - Waiting on
    - Mutex
    - File
    - Connection...
- Swapped out ...
  - Placed on disk temporarily





# MULTI-THREADED SYSTEMS



AARHUS  
UNIVERSITY

AARHUS UNIVERSITY SCHOOL OF ENGINEERING



# MULTI-THREADED SYSTEMS

- Advantages
- Disadvantages
  - Shared data
  - Starvation

# ADVANTAGES

- What are the advantages of multiple tasking a system?
  - Prioritization – the highest-priority task gets to run
  - Modularization – wrap concurrent activities in a task
  - Resource utilization – Don't spend CPU time waiting for I/O etc.
  - ...
- However, the use of multiple tasks in a system creates a number of problems for us
- We must know the problems to be able to identify and counter or avoid them

# DISADVANTAGES

## SHARED DATA

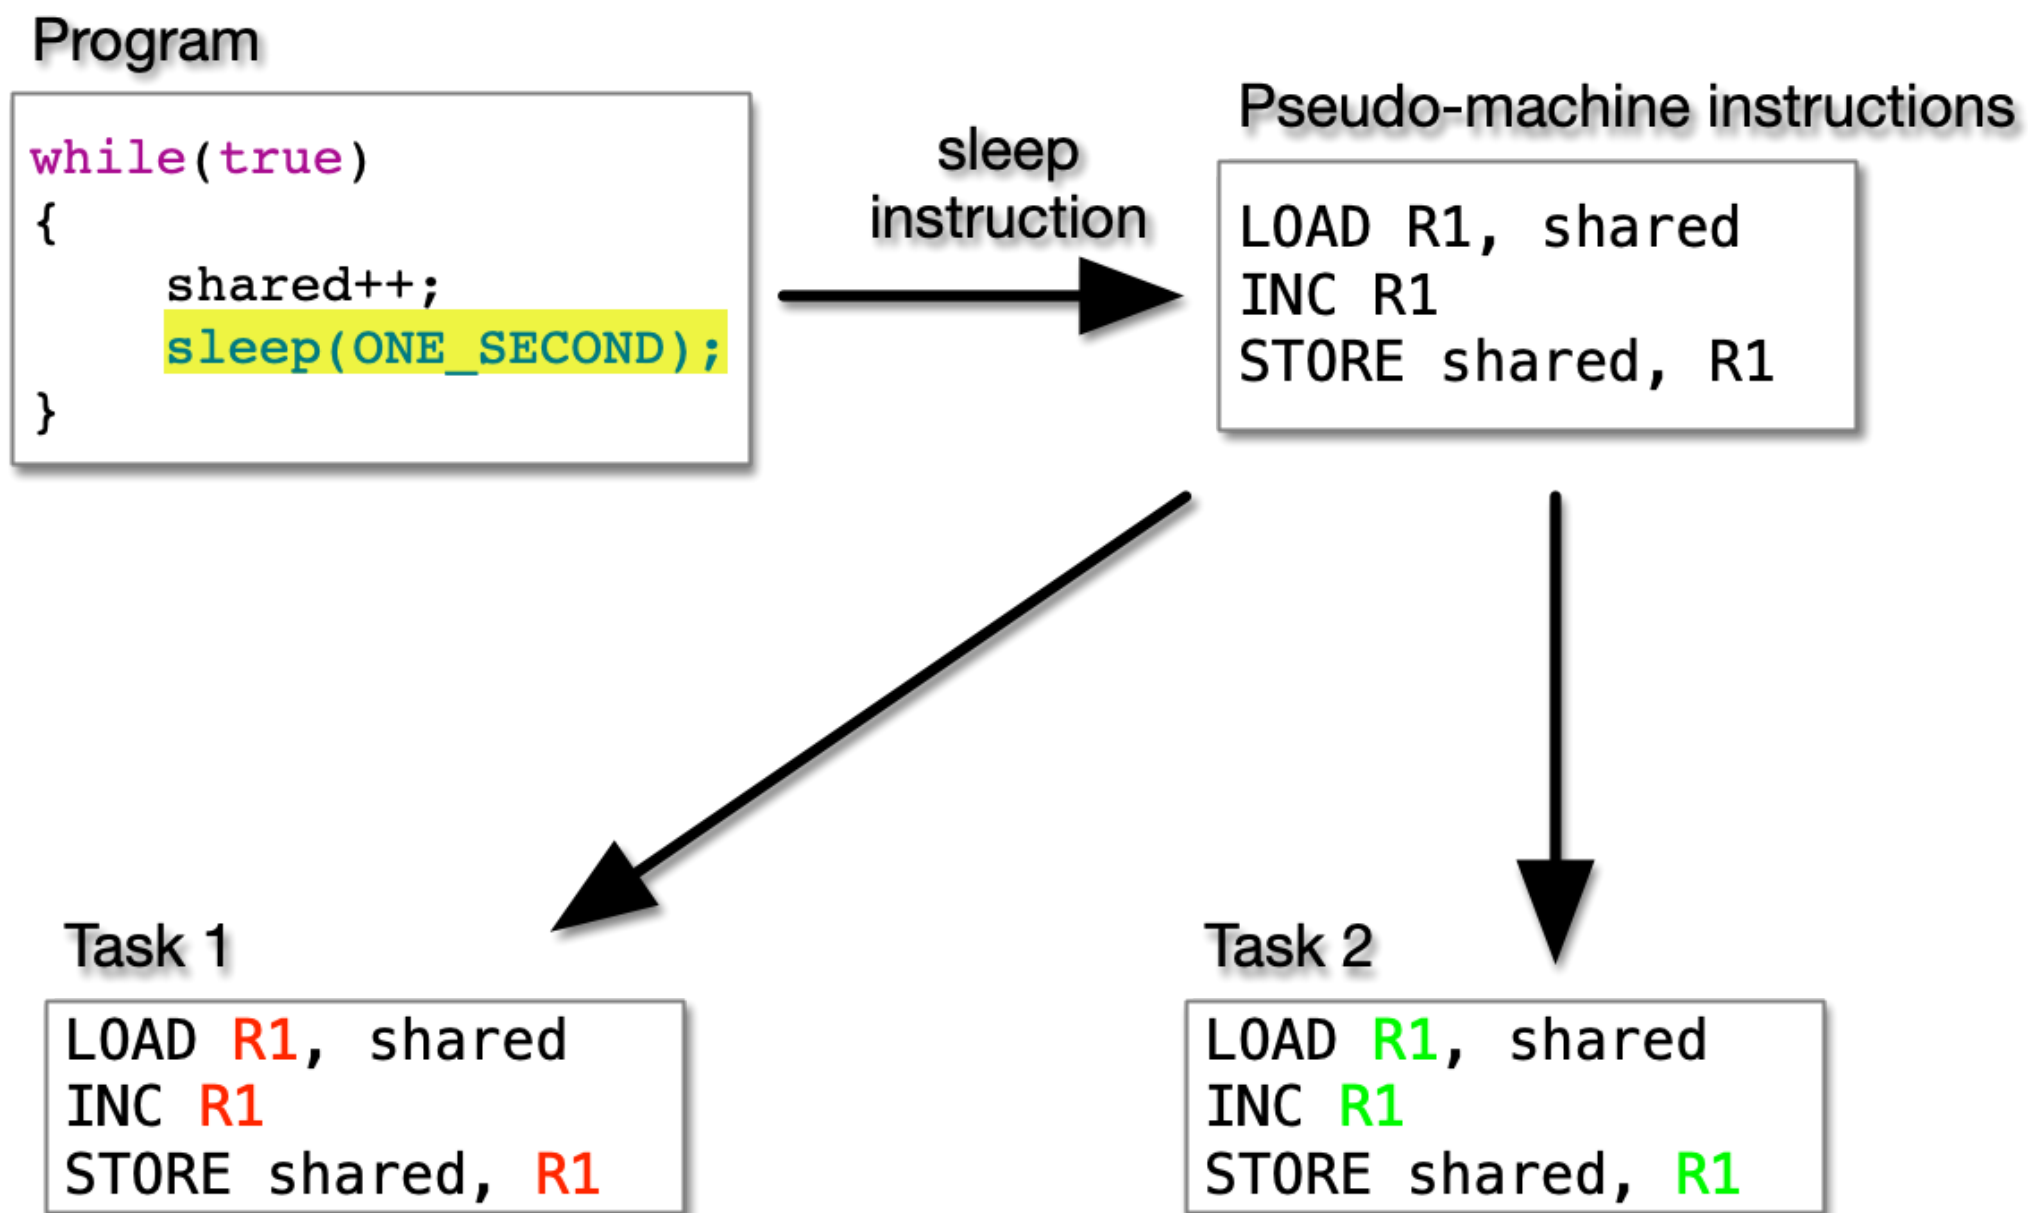
Consider the following code. What is the value of shared after 10 seconds?

```
01 unsigned int shared;
02 void taskfunc()
03 {
04     while(true)
05     {
06         shared++;           // Increment i, then wait
07         sleep(ONE_SECOND);  // 1 second
08     }
09 }
10
11 int main()
12 {
13     shared = 0;
14     createThread(taskFunc); // Start two identical threads
15     createThread(taskFunc); // that run the same function
16     while(true)
17         sleep();
18 }
```

# DISADVANTAGES

## SHARED DATA

Zooming in...



# DISADVANTAGES

## SHARED DATA

- Whats the difference between two execution scenarios?
- Who controls which scenario plays out?
- Which values can `shared` be?

Task 1

```
LOAD R1, shared
INC R1
STORE shared, R1
```

Task 2

```
LOAD R1, shared
INC R1
STORE shared, R1
```

Non-interleaved instructions

```
LOAD R1, shared // shared = 0
INC R1
STORE shared, R1 // shared = 1
LOAD R1, shared // shared = 1
INC R1
STORE shared, R1 // shared = 2
```

Interleaved instructions

```
LOAD R1, shared // shared = 0
LOAD R1, shared // shared = 0
INC R1
STORE shared, R1 // shared = 1
INC R1
STORE shared, R1 // shared = 1
```

# DISADVANTAGES

## SHARED DATA

```
01 struct Position
02 {
03     double x, y, z;
04 };
```

# DISADVANTAGES

## SHARED DATA

```
01 struct Position
02 {
03     double x, y, z;
04 };
```

```
01 void newPos(Position& pos, float x, float y, float z)
02 {
03     pos.x = x;
04     pos.y = y;
05     pos.z = z;
06 }
```



# DISADVANTAGES

## SHARED DATA

```
01 struct Position
02 {
03     double x, y, z;
04 };
```

```
01 void newPos(Position& pos, float x, float y, float z)
02 {
03     pos.x = x;
04     pos.y = y;
05     pos.z = z;
06 }
```

```
01 void printPos(Position& pos)
02 {
03     std::cout << "X: " << pos.x << std::endl;
04     std::cout << "Y: " << pos.y << std::endl;
05     std::cout << "Z: " << pos.z << std::endl;
06 }
```

# DISADVANTAGES

## SHARED DATA

```
01 struct Position
02 {
03     double x, y, z;
04 };
```

```
01 void newPos(Position& pos, float x, float y, float z)
02 {
03     pos.x = x;
04     pos.y = y;
05     pos.z = z;
06 }
```

```
01 void printPos(Position& pos)
02 {
03     std::cout << "X: " << pos.x << std::endl;
04     std::cout << "Y: " << pos.y << std::endl;
05     std::cout << "Z: " << pos.z << std::endl;
06 }
```

### GLOBAL COMMON VARIABLE

```
01 Position p { 10, 20, 30 };
```

#### TASK 1 (T1)

```
01 newPos(p, 11, 22, 33);
```

#### TASK 2 (T2)

```
01 printPos(p);
```

### NON-INTERLEAVED

```
01 T1 pos.x = x;
02 T1 pos.y = y;
03 T1 pos.z = z;
04
05 T2 std::cout << "X: " << pos.x << std::endl;
06 T2 std::cout << "Y: " << pos.y << std::endl;
07 T2 std::cout << "Z: " << pos.z << std::endl;
```

### INTERLEAVED

```
01 T1 pos.x = x;
02
03 T2 std::cout << "X: " << pos.x << std::endl; // X: 11
04 T2 std::cout << "Y: " << pos.y << std::endl; // Y: 20
05 T2 std::cout << "Z: " << pos.z << std::endl; // Z: 30
06
07 T1 pos.y = y;
08 T1 pos.z = z;
```

# DISADVANTAGES

## SHARED DATA

### NON-INTERLEAVED

```
01 T1 pos.x = x;  
02 T1 pos.y = y;  
03 T1 pos.z = z;  
04  
05 T2 std::cout << "X: " << pos.x << std::endl;  
06 T2 std::cout << "Y: " << pos.y << std::endl;  
07 T2 std::cout << "Z: " << pos.z << std::endl;
```

- Result
  - Ok, everything completes in *time*
  - Data is consistent

### INTERLEAVED

```
01 T1 pos.x = x;  
02  
03 T2 std::cout << "X: " << pos.x << std::endl; // X: 11  
04 T2 std::cout << "Y: " << pos.y << std::endl; // Y: 20  
05 T2 std::cout << "Z: " << pos.z << std::endl; // Z: 30  
06  
07 T1 pos.y = y;  
08 T1 pos.z = z;
```

- Result
  - Data skewed
    - Result should be have either
      - 10 & 20 & 30
      - 11 & 22 & 33
  - Unreliable

# DISADVANTAGES

## SHARED DATA


- The shared data problem is inherent in any preemptive multithreaded system
- Very cumbersome to find a good software solution to the problem
  - Peterson's solution: 2 interest flags, 1 will-wait flag
  - Does not scale
- We need a way to define critical sections of program
  - Sections in which the thread is guaranteed to be allowed to execute uninterrupted
- In a later lecture!

```
01 void taskfunc()  
02 {  
03     while(true)  
04     {  
05         enterCriticalSection();  
06         shared++;  
07         exitCriticalSection();  
08         sleep(ONE_SECOND);  
09     }  
10 }
```


# DISADVANTAGES

## STARVATION

- Consider a system with three threads, HP, MP and LP:
  - HP takes 5  $\mu$ s, must run once every 10  $\mu$ s
  - MP takes 3  $\mu$ s, must run once every 10  $\mu$ s
  - LP takes 1  $\mu$ s, must run once every 1000  $\mu$ s



SCHEDULE?




OK!


# DISADVANTAGES

## STARVATION


- Assume we add another MP thread, MP2
  - MP2 takes 2  $\mu$ s, may run once every 10  $\mu$ s



SCHEDULE?



**OK!**



**Starvation!**

# DISADVANTAGES

## STARVATION

- Starvation is an inherent problem in any priority-based system
- It occurs when the schedule is so tight that LP threads are never allowed to run because higher-priority threads “hog” the CPU
- Starvation can be very hard to predict and detect
  - Might only occur in very special situations

# PROGRAMMING WITH THREADS



AARHUS  
UNIVERSITY

AARHUS UNIVERSITY SCHOOL OF ENGINEERING





# PROGRAMMING WITH THREADS

- Summary
  - Threads
  - Processes
- Getting started



# THREADS - SUMMARY

- Threads are strands of execution – each process has at least one thread
  - AKA light-weight processes
  - Also called tasks or jobs
- Threads of the same process share memory space (e.g. global variables)
- Threads can harm each other
- We will often work with several threads in the same process



# PROCESSES - SUMMARY

- A process is an instance of a program that is being executed
  - Image of program (segment),
  - Stack, heap, registers, file descriptors, ...
- Processes have their own individual memory spaces
  - Process A cannot write in memory of process B – they are safe from each other
- Processes may only communicate through IPC mechanisms controlled by the OS.
- Processes may spawn other processes, which may execute the same or other programs
- A process may also spawn threads within its own memory space

# GETTING STARTED

- C++ does have the concept of threads built in
  - But we will not use it :-)
- The POSIX (Portable OS Interface for uniX) library is the most widely used threading library
  - Others include boost, Qt, ...
- The POSIX library has the thread type pthread which we will use.
  - Include `pthread.h`, link with library `pthread`



# GETTING STARTED

## PTHREAD FUNCTIONS

- Family of pthread functions you are to use...
  - `pthread_create()`
  - `pthread_join()`
  - `pthread_exit()`
  - `pthread_*` (and more)
- How they work is for next session