

Lecture 2

Why Does Software Fail?

- Background
 - What is Software Engineering?
 - What causes system failures?
 - What is the role of good engineering practice?
- Software failures vs. hardware failures
 - Example: Ariane-5 Flight 501
 - Example: Mars Exploration Rover (MER) Mission
- Lessons learned

1

Defining Software Engineering

- “Engineering...”
 - ... creates cost-effective solutions to practical problems by applying scientific knowledge to build **things** in the service of humankind”
- Software Engineering...
 - ... when the **things** contain software
- BUT:
 - Pure software is useless!
 - ... **software exists only as a part of a system**
 - Software is invisible, intangible, abstract
 - There are no physical laws underlying software behavior
 - There are no physical constraints on software complexity
 - Software never wears out
 - ... **traditional reliability measures do not apply**
 - Software can be replicated perfectly
 - ... **no manufacturing variability**

2

Failures and catastrophes

- (Hardware) system components often fail
 - Parts wear out
 - Wires and joints come loose
 - Components get used for things they were not designed for
 - Designs don't work the way they should
- Point failures typically do not lead to catastrophe
 - There are...
 - **backup systems**
 - **fault-tolerant designs**
 - **redundancy**
 - **certification using safety factors**
- Good engineering practice prevents accidents
 - Failure analysis
 - Reliability estimation
 - Checks and balances
- *But how does this work in Software Engineering?*

3

Ariane-5 Flight 501

- Background
 - The reusable launch vehicle of the European Space Agency (ESA)
 - Ariane-4 a major success
 - Ariane-5 developed for larger payloads
- Launched
 - 4 June 1996
- Mission
 - \$500 million payload to be delivered to orbit
- Fate
 - Veered off course during launch
 - Self-destructed approx. 40 sec after launch
- Cause
 - Unhandled floating-point exception in Ada code



4

Ariane-5 Events

- Locus of error
 - Platform alignment software (part of the Inertial Reference System, SRI)
 - This software only produces meaningful results prior to launch
 - Still operational 40 sec after launch
- Cause of error
 - Ada exception code raised and not handled
 - Converting 64-bit floating point to 16-bit signed integer for Horizontal Bias (BH)
 - Conversion worked fine in Ariane-4; overflow occurred in Ariane-5
 - Requirements state that computer shut down if unhandled exception occurs
- Launch+30s: SRIs fail
 - Backup SRI shuts down first
 - Active SRI shuts down 50ms later for same reason
- Launch+31s: On-Board Computer (OBC) receives data from active SRI
 - Diagnostic bit pattern interpreted as flight data (nonsensical input data)
 - OBC commands full nozzle deflections
 - Rocket veers off course
- Launch+33s: Launcher starts to disintegrate
 - Self-destruct triggers

5

Ariane-5 Why did this failure occur?

- Why was Platform Alignment still active after launch?
 - SRI software reused from Ariane-4
 - Ariane-5 was faster and used different timing parameters
- Why was there no exception handler?
 - To reduce the processor workload to below 80%
- Why wasn't the design modified for Ariane-5?
 - Not considered wise to change software that worked well for Ariane-4
- Why did the SRIs shut down?
 - They assumed faults are random hardware failures, hence should switch to backup
- Why was the error not caught in unit testing?
 - No trajectory data for Ariane-5 was provided in the requirements for SRIs
- Why was the error not caught in integration testing?
 - Full integration testing considered too expensive
 - SRIs were considered to be fully certified
 - Integration testing used simulations of the SRIs
- Why was the error not caught by inspection?
 - The implementation assumptions were not documented
- Why did the OBC diagnostic data as flight data?
 - They assumed this couldn't happen(!!)

6

Ariane-5 Summary

- Factors that lead to the Ariane-5 accident:
 - Didn't test all the specification assumptions
 - Insufficient test data
 - Lack of integration testing
 - Lack of expertise at inspections
 - Poor communication between teams
 - Commercial/management pressures took priority
 - Reused old code without checking newer assumptions
 - "Redundant" design not redundant

7

Mars Exploration Rover (MER)

- Background
 - The latest Mars mission of NASA
 - In spite of many problems, it's a major success
- Timeline
 - MER-A "Spirit"
 - Launched on Jun 10, 2003
 - Landed in Gusev Crater on Jan 4, 2004
 - MER-B "Opportunity"
 - Launched on Jul 7, 2003
 - Landed on Meridiani Planum on Jan 24, 2004
 - Designed to work for at least 90 days
 - Mars disappeared behind the Sun and out of the audio range on Sep 13, 2004
 - The rovers were still alive and running on Sep 13 (!)



8

MER Events

- Jan 21, 2004: Spirit stops communicating with the Earth
 - It reboots itself continuously
- Jan 26, 2004: The team discovers the cause
 - Corruption inside Spirit's on-board flash memory
 - The rovers are put into "deep sleep" until a patch is deployed
- February, 2004: The team beams up a software patch
 - Both rovers are up and running
- April, 2004: The team beams up another, non-critical software patch
- Locus of error
 - The DOS-like file allocation system, stored in the rover's flash memory
- Cause of error
 - Erasing files in DOS does not remove the file entry in directory, but marks it as "erased" and makes it available for further use
 - Creating many files at a time causes the directory to grow, without ever getting smaller (unless it's erased)
 - The flash memory was mirrored in the system's RAM, leading to overflows later on

9

MER Why did this failure occur?

- Why was a DOS-like file system used?
 - It's simple, fast, and it was successfully used on prior missions
 - Extreme radiation conditions on Mars require a slow CPU clock
 - Software cannot be too complex
- Why was the file system overloaded?
 - A large number of tiny files were created during data acquisition
- Why did the operating system crash?
 - The data stored in flash memory (including file system) was mirrored in the system's RAM
 - The RAM was nearly full, and further dynamic memory allocations resulted in a crash
- Why was the error not caught in prior testing?
 - NASA's tests only allowed for the addition of a small number of data files
- How was the error fixed?
 - The dynamic allocation feature was disabled
 - The flash memory was erased
 - A file system monitoring utility was installed

10

MER Summary

- Factors that lead to the flash memory problems:
 - Didn't test all specification assumptions
 - Insufficient test data
 - Lack of integration testing
 - Poor communication between teams
 - Reused code without checking assumptions
 - Others
 - ... to be disclosed by NASA (?)

11

Lessons learned

- Failures can usually be traced to a single root cause
 - ... but good engineering practice should prevent system failures
 - The real problems are failures of:
 - **testing and inspection process**
 - **problem reporting and tracking**
 - **lack of expertise**
 - **inadequate resources**
 - **etc.**
 - In most cases, it takes a failure of both engineering practice and management to cause a system failure
- Reliable software depends not on writing flawless code, but on how good are we at:
 - Communication (sharing information between teams)
 - Management (of resources and skills)
 - Verification and validation
 - Risk identification and tracking
 - Questioning assumptions

12

Readings

- Hans van Vliet's book, Chapter 1
 - Read all of it
 - Pay attention to the code of ethics
- Ariane-5
 - Information about ESA's launches
 - <http://www.esa.int/export/esaLA/launchers.html>
 - Flight 501 inquiry report & Press release:
 - <http://www.esrin.esa.it/htdocs/tidc/Press/Press96/press33.html>
- Mars Exploration Rover (MER) Mission
 - Information about NASA's MER Mission
 - <http://marsrovers.jpl.nasa.gov/home/>
 - DOS glitch nearly killed Mars Rover
 - <http://www.extremetech.com/article2/0,1558,1638764,00.asp>