



Debugging:

how I learned to stop hacking and love my
tools





What You Won't Learn

- how to solve all your problems
- the one true way to avoid all bugs
- Many platform specific debugging techniques





Documentation

- learn where the documentation for your system is and use it
- when in doubt, look it up!
 - Internet is your friend, but many clueless people so must filter. Use definitive sources if you can
 - man pages/Info system for Unix systems
 - MSDN for windows
 - javadoc for Java libraries
 - /Developer/Documentation for Mac OS X





Debugging

- Debugging is a black art. Some things to go over, though, so they'll be concrete in our brains:
 - relation to testing
 - why debugging is hard
 - types of bugs
 - process
 - techniques
 - tools
 - avoiding bugs





Debugging and testing

- Testing and debugging go together like peas in a pod:
 - Testing finds errors; debugging localizes and repairs them.
 - Together these form the “testing/debugging cycle”: we test, then debug, then repeat.
 - Any debugging should be followed by a reapplication of *all* relevant tests, particularly regression tests. This avoids (reduces) the introduction of new bugs when debugging.
 - Testing and debugging need not be done by the same people (and often should not be).





Why debugging is hard

- There may be no obvious relationship between the external manifestation(s) of an error and its internal cause(s).
- Symptom and cause may be in remote parts of the program.
- Changes (new features, bug fixes) in program may mask (or modify) bugs.
- Symptom may be due to human mistake or misunderstanding that is difficult to trace.
- Bug may be triggered by rare or difficult to reproduce input sequence, program timing (threads) or other external causes.
- Bug may depend on other software/system state, things others did to you systems weeks/months ago.





Designing for Debug/Test

- when you write code think about how you are going to test/debug it
 - lack of thought *always* translates into bugs
- write test cases when you write your code
- if something should be true `assert()` it
- create functions to help visualize your data
- design for testing/debugging from the start
- test early, test often
- test at abstraction boundaries





Fault Injection

- many bugs only happen in the uncommon case
- make this case more common having switches that cause routines to fail
 - file open, file write, memory allocation, are all good candidates
- Have “test drivers” which test with the uncommon data. If deeply buried, test with a debugger script





Finding and Fixing Bugs

- in order to create quality software you need to find your bugs
 - testing
 - user reports
- the best bugs are those that are always reproducible





Types of bugs

- Types of bugs (gotta love em):
 - Compile time: syntax, spelling, static type mismatch.
 - Usually caught with compiler
 - Design: flawed algorithm.
 - Incorrect outputs
 - Program logic (if/else, loop termination, select case, etc).
 - Incorrect outputs
 - Memory nonsense: null pointers, array bounds, bad types, leaks.
 - Runtime exceptions
 - Interface errors between modules, threads, programs (in particular, with shared resources: sockets, files, memory, etc).
 - Runtime Exceptions
 - Off-nominal conditions: failure of some part of software of underlying machinery (network, etc).
 - Incomplete functionality
 - Deadlocks: multiple processes fighting for a resource.
 - Freeze ups, never ending processes





The ideal debugging process

- A debugging algorithm for software engineers:
 - Identify test case(s) that reliably show existence of fault (when possible)
 - Isolate problem to small fragment(s) of program
 - Correlate incorrect behavior with program logic/code error
 - Change the program (and check for other parts of program where same or similar program logic may also occur)
 - Regression test to verify that the error has really been removed - without inserting new errors
 - Update documentation when appropriate

(Not all these steps need be done by the same person!)



General Advice

- try to understand as much of what is happening as possible
- “it compiles” is **NOT** the same as “it works”
- when in doubt, ask. Then test the answer!
- Error messages are generally just a vague hint and can be misleading.
- Don’t always trust the “comments/documents”, they can be out-of-date.





What is a Debugger?

“A software tool that is used to detect the source of program or script errors, by performing step-by-step execution of application code and viewing the content of code variables.”

-MSDN





What is a Debugger? (con't)

- A debugger is *not an IDE*
 - Though the two can be integrated, they are separate entities.
- A debugger loads in a program (compiled executable, or interpreted source code) and allows the user to trace through the execution.
- Debuggers typically can do disassembly, stack traces, expression watches, and more.





Why use a Debugger?

- No need for precognition of what the error might be.
- Flexible
 - Allows for “live” error checking – no need to re-write and re-compile when you realize a certain type of error may be occurring
 - Dynamic
 - Can view the entire relevant scope





Why people don't *use a Debugger*?

- With simple errors, may not want to bother with starting up the debugger environment.
 - Obvious error
 - Simple to check using prints/asserts
- Hard-to-use debugger environment
- Error occurs in optimized code
- Changes execution of program (error doesn't occur while running debugger)





Debugging techniques, 1

- Execution tracing
 - running the program
 - print
 - trace utilities
 - single stepping in debugger
 - hand simulation





Debugging techniques, 2

- Interface checking
 - check procedure parameter number/type (if not enforced by compiler) and value
 - *defensive programming*: check inputs/results from other modules
 - documents assumptions about caller/callee relationships in modules, communication protocols, etc
- Assertions: include range constraints or other information with data.
- Skipping code: comment out suspect code, then check if error remains.



Other Functions of a Debugger

- Disassembly (in context and with live data!)
- Execution Tracing/Stack tracing
- Symbol watches





Disassembly

- Most basic form of debugging
- Translating machine code into assembly instructions that are more easily understood by the user.
- Typically implementable as a simple lookup table
- No higher-level information (variable names, etc.)
- Relatively easy to implement.





Execution Tracing

- Follows the program through the execution. Users can step through line-by-line, or use breakpoints.
- Typically allows for “watches” on – registers, memory locations, symbols
- Allows for tracing up the stack of runtime errors (back traces)
- Allows user to trace the causes of unexpected behavior and fix them





Symbol Information

- Problem – a compiler/assembler translates variable names and other symbols into internally consistent memory addresses
- How does a debugger know which location is denoted by a particular symbol?
- We need a “debug” executable.





Debug vs. Release Builds

- Debug builds usually are *not optimized*
- Debug executables contain:
 - program's symbol tables
 - location of the source file
 - line number tags for assembly instructions.
- GCC/GDB allows debugging of optimized code.





Bug hunting with print

- Weak form of debugging, but still common
- How bug hunting with print can be made more useful:
 - print variables other than just those you think suspect.
 - print valuable statements (not just “hi\n”).
 - use `exit()` to concentrate on a part of a program.
 - move print through a through program to track down a bug.



Debugging with print (continued)

- Building debugging with print into a program (more common/valuable):
 - print messages, variables/test results in useful places throughout program.
 - use a ‘debug’ or ‘debug_level’ global flag to turn debugging messages on or off, or change “levels”
 - possibly use a source file preprocessor (#ifdef) to insert/remove debug statements.
 - Often part of “regression testing” so automated scripts can test output of many things at once.



Finding Reproducible Bugs

- a bug happens when there is a mismatch between what you (someone) *think* is happening and what is *actually* happening
- confirm things you believe are true
- narrow down the causes one by one
- make sure you understand your program state
- keep a log of events and assumptions





Finding Reproducible Bugs

- try explaining what should be happening
 - Verbalization/writing often clarifies muddled thoughts
- have a friend do a quick sanity check
- don't randomly change things, your actions should have a purpose.
 - If you are not willing to check it into CVS with a log that your boss may read, then you are not ready to make that change to the code.
 - Think it through first, both locally and globally.





(semi) irreproducible bugs

- sometimes undesired behavior only happens sporadically
- tracking down these heisenbugs is hard
- the error could be at any level
 - Circuits (e.g. bad ram chip at high memory address)
 - compiler
 - OS
 - Linker
 - Irreproducible external “data” and timing





Finding HeisenBugs

- Use good tools like Purify. Most common “Heisenbugs” are memory or thread related.
- try to make the bug reproducible by switching platforms/libraries/compilers
- insert checks for invariants and have the program stop everything when one is violated
- verify each layer with small, simple tests
- find the smallest system which demonstrates the bug
- Test with “canned data”, replayed over net if needed.



Timing and Threading Bugs

- ensure the functionality works for a single thread
- if adding a `printf()` removes the bug it is almost certainly a timing/threading bug or a trashed memory bug
- try using coarse grained locking to narrow down the objects involved
- try keeping an event (transaction) log





Memory Bugs and Buffer Overflow

- Trashing stack/heap causes often difficult to find bugs.
- Manifestation can be far from actual bug.
 - “Free list” information generally stored just after a “malloced” chunk of data. Overwriting may not cause problem until data is “freed”, or until something else does a malloc after the free.
 - Stack variables, overwriting past end, changes other variables, sometimes return address. (Buffer overflow)
 - Bad “string” ops notorious, using input data can also be problematic.





An example....

```
void myinit(int startindex, int startvalue, int length, int* vect){
    int i;
    for(i=startindex; i< startindex+length; i++)
        *vect++ = startvalue++;
}

void whattheheck(){    printf("How did I ever get here????\n");    exit(2); }

int main(int argc, char**argv){
    float d;
    int a,b[10],c, i, start,end;
    if(argc != 3) {printf("Usage:%s start, end\n",argv[0]);exit(-1); }
    start=atoi(argv[1]);    end=atoi(argv[2]);    /* bad style but shorter */
    a=0;    c=0;    d=3.14159;    /* bad style but shorter */
    printf("Initially a %d, c %d, d %f, start %d, end %d\n",a,c,d, start,end);
    myinit(start,start,end,b+start);
    printf("finally a %d, c %d, d %f start %d, end %d \n",a,c,d, start, end);
    if(end>10) b[end-1]=134513723;
    return 0;
}
```





Debugging Techniques

- methodology is key
- knowing about lots of debugging tools helps
- the most critical tool in your arsenal is your brain
- second most important tool is a debugger
 - Core dumps are your friends.. Learn how to use them.
- Memory debugging tools third most important
- Profiler fourth





Tools

- Tracing programs:
 - strace, truss (print out system calls),
 - /bin/bash -X to get a shell script to say what its doing
- Command-line debuggers :
 - gdb (C, C++), jdb (java), “perl -d”
- Random stuff:
 - electric fence or malloc_debug, mtrace, etc (a specialized malloc() for finding memory leaks in C/C++)
 - purify (Part of Rational Suite, a really good memory debugging tools for C/C++ programs)





Debuggers

- allow programs to inspect the state of a running program
- become more important when debugging graphical or complex applications
- jdb and gjdb for java
- gdb for C/C++ programs on unix/Max
- MSVS for Win32
- kdb for the linux kernel





Using gdb

- Compile debugging information into your program:

```
gcc -g <program>
```

- Read the manual:

- `man gdb`

- in emacs (tex-info):

- `http://www.cslab.vt.edu/manuals/gdb/gdb_toc.html`

- `M-x help <return> i <return>, arrow to GDB, <return>`

- online:





`gdb` commands

`run/continue/next/step`: start the program, continue running until break, next line (don't step into subroutines), next line (step into subroutines).

`break <name>`: set a break point on the named subroutine. Debugger will halt program execution when subroutine called.

`backtrace`: print a stack trace.

`print <expr>`: execute expression in the current program state, and print results (expression may contain assignments, function calls).

`help`: print the help menu. `help <subject>` prints a subject menu.

`quit`: exit `gdb`.



General GDB

- easiest to use inside of emacs (M-x gdb)
- run starts the program
- set args controls the arguments to the program
- breakpoints control where the debugger stops the program (set with C-x space)
- next moves one step forward
- step moves one step forward and also traverses function calls
- continue runs the program until the next breakpoint





General GDB

- p prints a value (can be formatted)
 - /x hex
 - /o octal
 - /t binary (t is for two)
 - /f float
 - /u unsigned decimal





General GDB

- `bt` shows the current backtrace (also where)
- `up/down` move up down the stack
- `f #` lets you switch to frame # in the current backtrace
- `set var=value`
- `call` allows call of a function (the syntax can get funky)
- `jump` (dump to a particular line of code)
- `Thread #` lets you switch to a particular thread



Advanced GDB

- watchpoints let you check if an expression changes
- catchpoints let you know when interesting things like exec calls or library loads happen
- x lets you inspect the contents of memory
- Watchpoints can be quite slow, best to combine with breakpoints.





Advanced GDB

- gcc 3.1 and up provides macro information to gdb if you specify the options -gdwarf2 and -g3 on the command line
- you can debug and already running process with the attach *pid* command
- you can apply a GDB command to all threads with thread apply all
- GDB can be used to debug remote embedded systems (gdbserver, etc..)





The example in gdb

```
void myinit(int startindex, int startvalue, int length, int* vect){
    int i;
    for(i=startindex; i< startindex+length; i++)
        *vect++ = startvalue++;
}

void whattheheck(){    printf("How did I ever get here????\n");    exit(2); }

int main(int argc, char**argv){
    float d;
    int a,b[10],c, i, start,end;
    if(argc != 3) {printf("Usage:%s start, end\n",argv[0]);exit(-1); }
    start=atoi(argv[1]);    end=atoi(argv[2]);    /* bad style but shorter */
    a=0;    c=0;    d=3.14159;    /* bad style but shorter */
    printf("Initially a %d, c %d, d %f, start %d, end %d\n",a,c,d, start,end);
    myinit(start,start,end,b+start);
    printf("finally a %d, c %d, d %f start %d, end %d \n",a,c,d, start, end);
    if(end>10) b[end-1]=134513723;
    return 0;
}
```





JDB

- Compile your source file with **-g** option
 - Produce debugging information
 - For example
javac -g rsdimu.java
- To run jdb,
 - Type “jdb”, use run command to load an executable
 - **run <class name>**
 - For example, **run rsdimu**
 - OR type “jdb <class name>”, then
 - Type “run” to execute the program
 - Type “quit” to exit jdb





Breakpoint

- Make your program stops whenever a certain point in the program is reached

- For example

- Make a breakpoint in line 6

stop at Demo:6

- Program stops before execute line 6
- Allow you to examine code, variables, etc.

```
1 public class Demo
2 {
3     public static void main(...)
4     {
5         System.out.println
6         ("A\n");
7         System.out.println
8         ("B\n");
9         System.out.println
```



Breakpoint

- Add breakpoint
 - stop at MyClass:<line num>
 - stop in java.lang.String.length
 - stop in MyClass.<method name>
- Delete breakpoint
 - clear (clear all breakpoints)
 - clear <breakpoint>
 - e.g. clear MyClass:22





Step

- Execute one source line, then stop and return to JDB
- Example

```
public void func()
{
    System.out.println
("A\n");
    System.out.println
("B\n");
    System.out.println
```

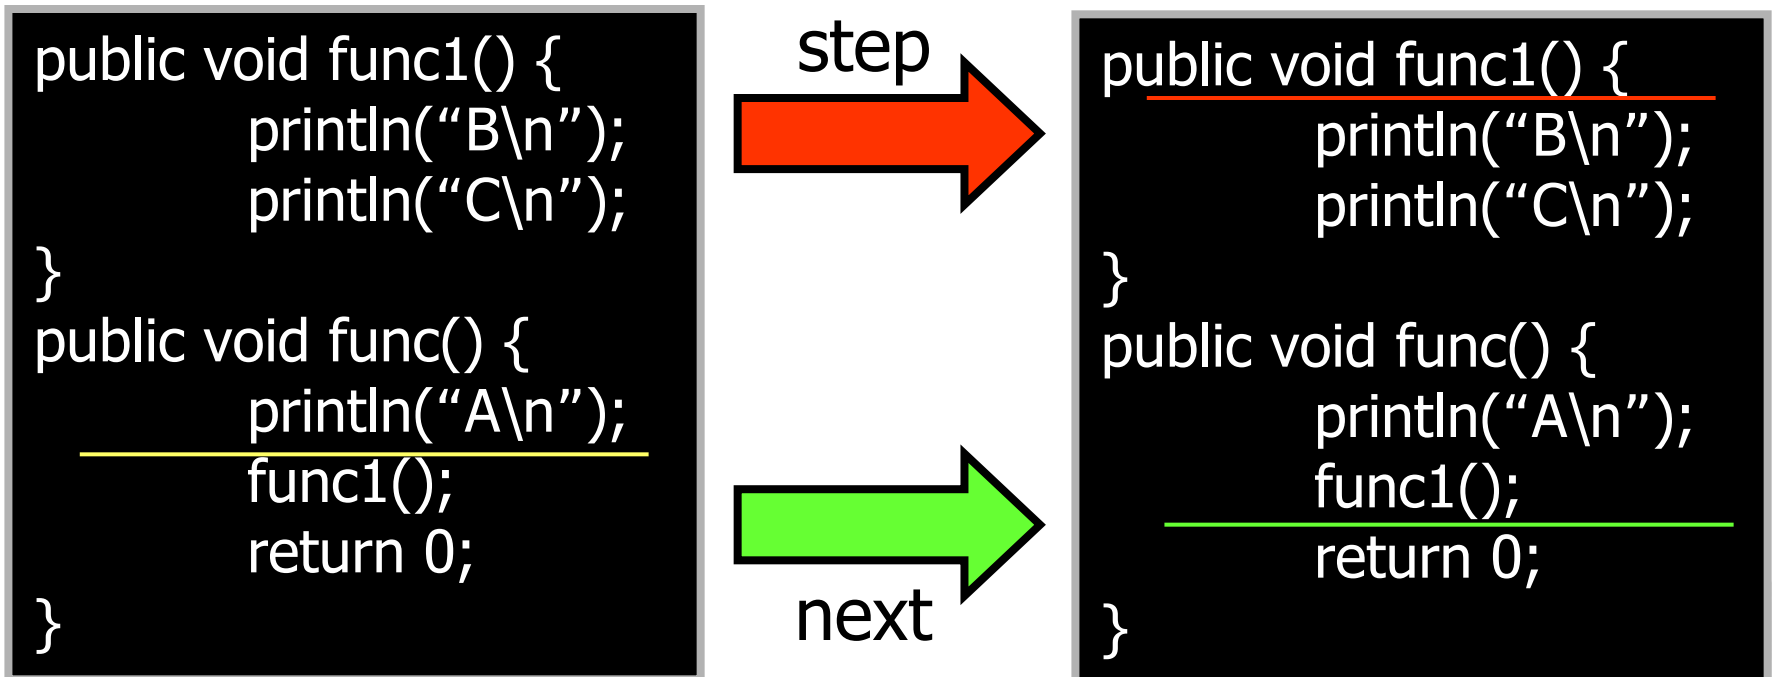
step



```
public void func()
{
    System.out.println("A\n");
    System.out.println("B\n");
    System.out.println("C\n");
    return 0;
}
```

Next

- Similar to step, but treat function call as one source line
- Example





Cont

- Resume continuous execution of the program until either one of the followings
 - Next breakpoint
 - End of program





Print

■ Print

- Display the value of an expression

- **print expression**

- **print MyClass.myStaticField**
- **print i + j + k**
- **print myObj.myMethod()** (*if myMethod returns a non-null*)
- **print new java.lang.String("Hello").length()**

■ Dump

- Display all the content of an object

- **dump <object>**





Analysis Tools

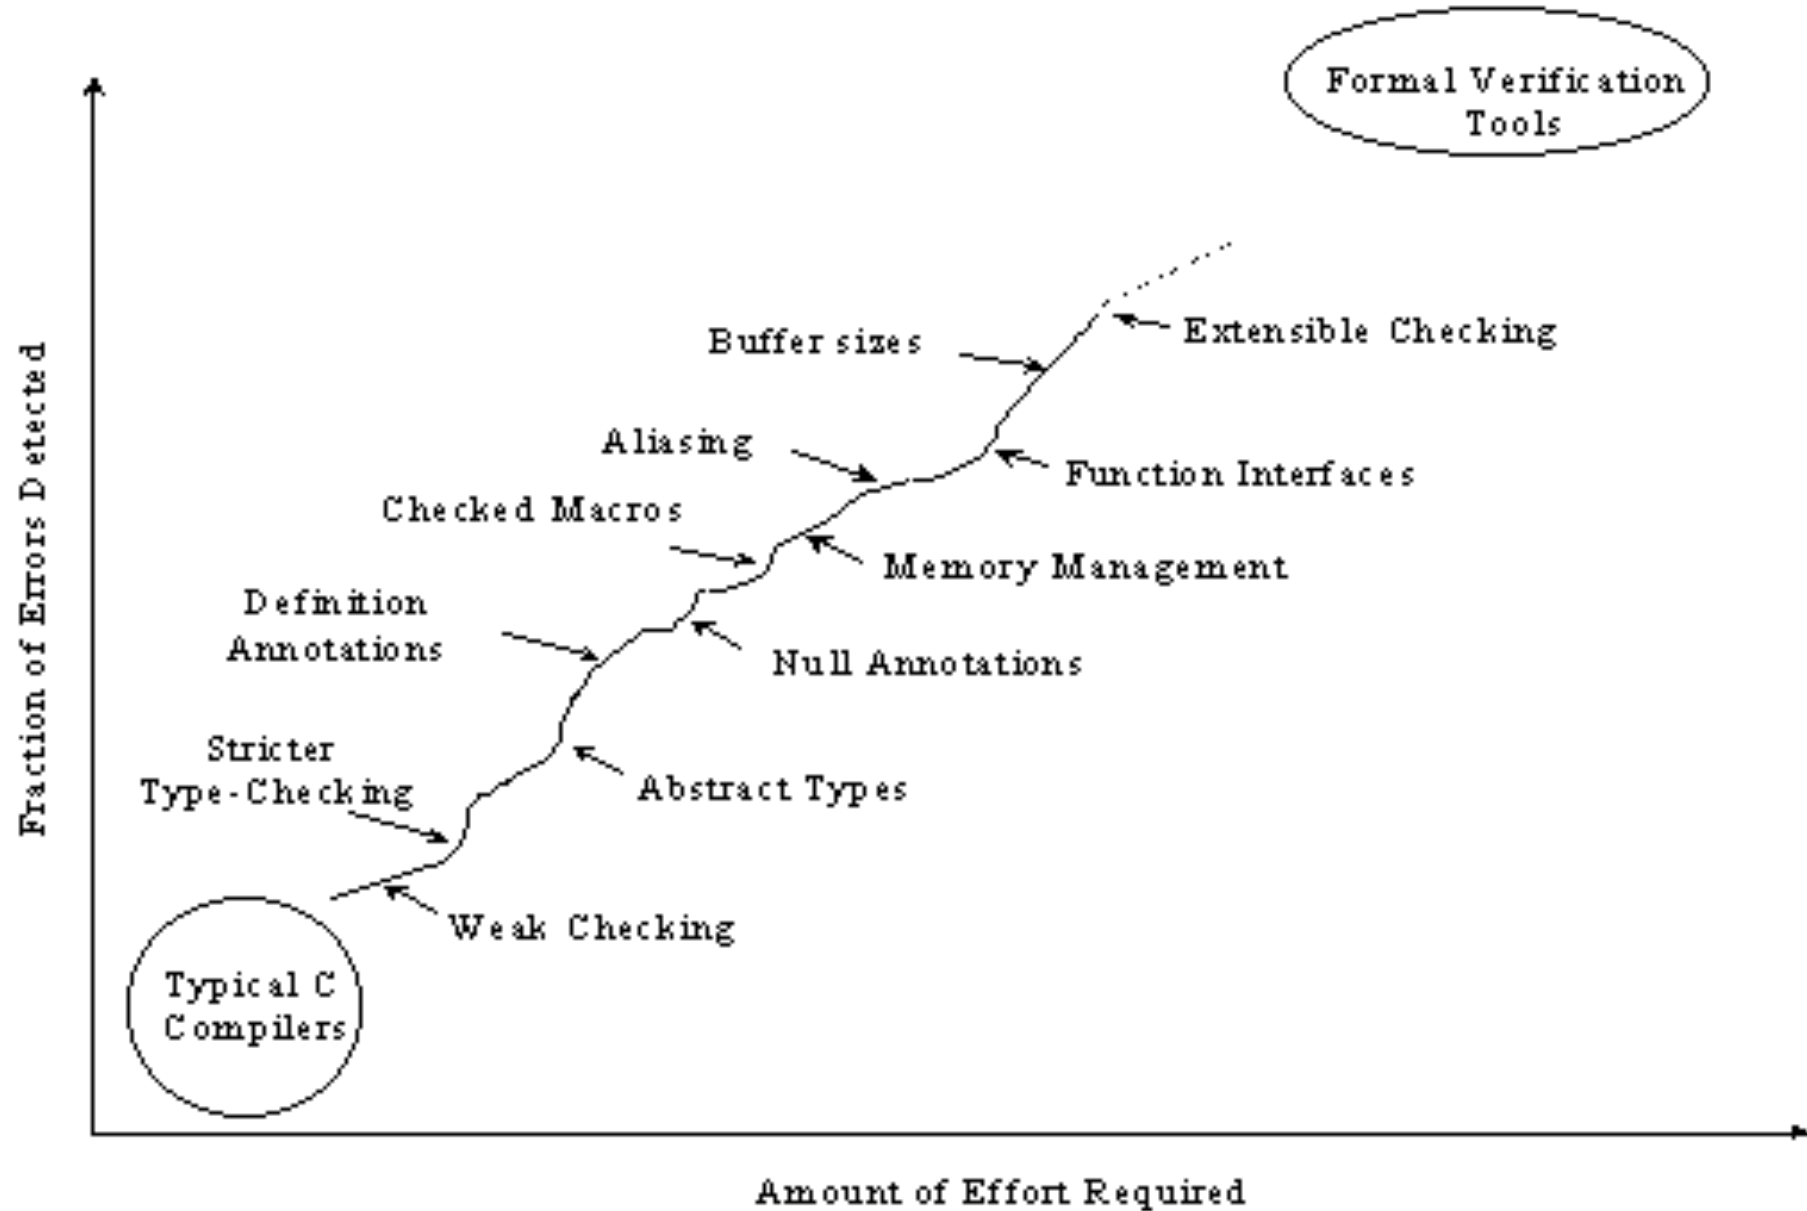




Purpose of Analysis Tools

- Need for a feasible method to catch bugs in large projects. Formal verification techniques require unreasonable effort on large projects.
- Augment traditional debugging techniques without adding unreasonable burden to the development process.







Two Types of Analysis Tools

- Static Analysis
- Run-time (dynamic) Analysis





Static Analysis

- Examine a program for bugs without running the program.
- Examples:
 - Splint (www.splint.org),
 - PolySpace C Verifier (www.polyspace.com).





Splint

- Open Source Static Analysis Tool developed at U.Va by Professor Dave Evans.
- Based on Lint.





Errors Splint will detect

- Dereferencing a possibly null pointer.
- Using possibly undefined storage or returning storage that is not properly defined.
- Type mismatches, with greater precision and flexibility than provided by C compilers.
- Violations of information hiding.
- Memory management errors including uses of dangling references and memory leaks.
- Dangerous aliasing.





Errors Splint will detect continued...

- Modifications and global variable uses that are inconsistent with specified interfaces.
- Problematic control flow such as likely infinite loops.
- Buffer overflow vulnerabilities.
- Dangerous macro initializations and invocations.
- Violations of customized naming conventions.





What's wrong with this code?

```
void strcpy(char* str1, char* str2)
{
    while (str2 != 0)
    {
        *str1 = *str2;
        str1++;
        str2++;
    }
    str1 = 0; //null terminate the string
}
```





What happens to the stack?

```
void foo()
{
    char buff1[20]; char buff2[40];
    ...
    //put some data into buff2
    strcpy(buff1, buff2);
}
```





Secure Programming

- Exploitable bugs such as buffer overflows in software are the most costly bugs.
- Incredibly frequent because they are so hard to catch.
- Analysis tools play a big part in finding and fixing security holes in software.





How does Splint deal with false positives?

- Splint supports annotations to the code that document assumptions the programmer makes about a given piece of code
- These annotations help limit false positives.





Run-time Analysis

- Many bugs cannot be determined at compile time. Run-time tools required to find these bugs.
- Run-time analysis tools work at run-time instead of compile time.
- Example – Purify (www.rational.com).





Purify

- Purify modifies object files at link time.
- After execution, Purify will report bugs such as memory leaks and null dereferences.





Purify continued...

- From the purify manual: “Purify checks every memory access operation, pinpointing where errors occur and providing diagnostic information to help you analyze why the errors occur.”





Types of errors found by Purify

- Reading or writing beyond the bounds of an array.
- Using un-initialized memory.
- Reading or writing freed memory.
- Reading or writing beyond the stack pointer.
- Reading or writing through null pointers.
- Leaking memory and file descriptors.
- Using a pointer whose memory location was just deallocated





Static vs. Run-time Analysis

- Probably good to use both.
- Run-time analysis has fewer false positives, but usually requires that a test harness test all possible control flow paths.





Cons of Analysis Tools

- Add time and effort to the development process.
- Lots of false positives.
- No guarantee of catching every bug.
- However, in a commercial situation, probably worth your time to use these tools.





Other Tools

- Mallocdebug and debugmalloc libraries
- valgrind is a purify-like tool which can really help track down memory corruption (linux only)
- MemProf for memory profiling and leak detection (linux only) www.gnome.org/projects/memprof
- electric fence is a library which helps you find memory errors
- c++filt demangles a mangled c++ name





Memory Leaks

- Memory bugs
 - Memory corruption: dangling refs, buffer overflows
 - Memory leaks
 - Lost objects: unreachable but not freed
 - Useless objects: reachable but not used again





Memory Leaks

- Memory bugs
 - Memory corruption: dangling refs, buffer overflows
 - Memory leaks
 - Lost objects: unreachable but not freed
 - Useless objects: reachable but not used again

Managed Languages

- 80% of new software in Java or C# by 2010
[Gartner] (personally TB does not believe it..)
- Type safety & GC eliminate many bugs



Memory Leaks

- Memory bugs
 - ~~Memory corruption: dangling refs, buffer overflows~~
 - Memory leaks
 - ~~Lost objects: unreachable but not freed~~
 - Useless objects: “reachable” but not used again

Managed Languages

- 80% of new software in Java or C# by 2010
[Gartner]
- Type safety & GC eliminate many bugs



Memory Leaks

Leaks occur in practice in managed languages

[Cork, JRockit, JProbe, LeakBot, .NET Memory Profiler]

– Memory leaks

- ~~Lost objects: unreachable but not freed~~
- Useless objects: reachable but not used again

Managed Languages

- 80% of new software in Java or C# by 2010
[Gartner]
- Type safety & GC eliminate many bugs



Other linux Tools

- strace / truss / ktrace let you know what system calls a process is making
- Data Display Debugger (DDD) is good for visualizing your program data www.gnu.org/software/ddd
- gcov lets you see which parts of your code are getting executed
- Profilers (gprof) to see where you are spending “time” which can help with performance logic bugs



Code beautifier

- Improve indentation of your source code for better readability
- source code beautifier in UNIX/Cygwin
 - Indent
 - M-x indent-region in emacs
- Make sure the code beautifier does not change how your code works after beautification!





Avoiding bugs in the first place

- Coding style: use clear, consistent style and useful naming standards.
- Document everything, from architecture and interface specification documents to comments on code lines.
- Hold code reviews.
- Program defensively.
- Use/implement exception handling liberally; think constantly about anomalous conditions.
- Be suspicious of cut/paste.
- Consider using an integrated development environment (IDE) with dynamic syntax checking






Code reviews

- Primary programmer(s) for some piece of code presents and explains that code, line by line.
- Audience of programmers experienced in language, code's general domain. Audience may also contain designers, testers, customers and others less versed in code but concerned with quality and consistency.
- Review is a dialogue: audience pushes presenters to reevaluate and rationalize their implementation decisions.
- Extremely useful: reviews often turn up outright errors, inconsistencies, inefficiencies and unconsidered exceptional conditions.
- Also useful in familiarizing a project team with a member's code.





War Stories





Nasty Problems

- overwriting return address on the stack
- overwriting virtual function tables
- compiler bugs - rare but my students and I have a talent for finding them
 - try -O, -O2, -Os
- wrong version of a shared library gets loaded
 - make the runtime linker be verbose





Nasty Problems

- Difference in MS “debug/development” stacks vs performance/runtime stacks
- OS / library bugs
 - create small examples
 - verify preconditions
 - check that postconditions fail
- static initialization problems with C++
- Processor bugs





some fun (simple) bugs

- In java (or whatever):

```
public void foo(int p, int q) {  
    int x = p;  
    int y = p;  
}
```

- In perl:

```
$foo = $cgi->param('foo');  
if (!$foo) {  
    webDie ("missing parameter foo!");  
}
```

- In C:

```
char *glue(char *left, char sep, char *right) {  
    char *buf = malloc(sizeof(char) *  
                        (strlen(left) + 1 strlen(right)));  
    sprintf(buf, "%s%c%s", left, sep, right);  
    return buf;  
}
```





Tricks

- write a custom `assert()`
- stub routines on which to break
- dump to a file instead of standard out
- make `rand()` deterministic by controlling the seed
- fight memory corruption with tools that use sentinels, etc. (and if no tools do it yourself)





The Future of Debugging

- better debuggers and programs to help you visualize your programs state
- simple model checkers
- programs keep getting bigger, finding bugs is going to get harder!
- Parallel/distributed debuggers as we move to more parallel/distributed systems.

